



BroadView

BroadView™ Instrumentation Agent Specification

Specification

Software Release 3.2

Broadcom, the pulse logo, Connecting everything, Avago Technologies, Avago, and the A logo are among the trademarks of Broadcom and/or its affiliates in the United States, certain other countries, and/or the EU.

Copyright © 2017–2018 Broadcom. All Rights Reserved.

The term “Broadcom” refers to Broadcom Inc. and/or its subsidiaries. For more information, please visit www.broadcom.com.

Broadcom reserves the right to make changes without further notice to any products or data herein to improve reliability, function, or design. Information furnished by Broadcom is believed to be accurate and reliable. However, Broadcom does not assume any liability arising out of the application or use of this information, nor the application or use of any product or circuit described herein, neither does it convey any license under its patent rights nor the rights of others.

Table of Contents

Chapter 1: Introduction	7
1.1 Overview	7
1.2 Supported Features by Silicon	8
Chapter 2: Architecture	9
2.1 Architecture Blocks	10
2.1.1 Northbound Plug-in Block	10
2.1.1.1 Reference Implementation.....	10
2.1.2 Infrastructure Block	11
2.1.2.1 Reference Implementation.....	11
2.1.3 Southbound Plug-in Block.....	11
2.1.3.1 Reference Implementation.....	11
2.1.4 Instrumentation Features Block	12
2.2 Execution in Linux	12
2.3 Using the Agent with a NOS.....	13
2.3.1 SB Plug-ins	13
2.3.2 NB Plug-ins	13
Chapter 3: Communication	14
3.1 REST.....	14
3.2 JSON-RPC Messaging	14
3.2.1 Request Response Messages	14
3.2.2 Time Stamps.....	15
3.2.3 Version.....	15
3.3 Multi-ASIC Platforms	16
3.4 Error Reporting	17
Chapter 4: Generic Agent Management	18
4.1 Overview	18
4.2 Heartbeat Messaging	18
4.3 Switch Properties	19
4.4 Configurability	20
4.5 API	21
4.5.1 Overview	21
4.5.2 Configuration API.....	21
4.5.2.1 configure-system-feature	21
4.5.2.2 save-configuration	22
4.5.2.3 restore-configuration.....	22
4.5.2.4 cancel-request	23
4.5.3 Status/Reporting API	24

4.5.3.1 get-configuration	24
4.5.3.2 get-port-queue-map	26
4.5.3.3 get-system-feature	35
4.5.3.4 get-switch-properties	35
Chapter 5: Buffer Statistics and Tracking	37
5.1 Overview	37
5.2 Buffer Utilization Statistics and Thresholds	37
5.3 Configurability	38
5.4 Complete and Incremental Reports	39
5.5 Queues	39
5.6 Deciphering Reports	40
5.6.1 Example 1	40
5.6.2 Example 2	41
5.6.3 Example 3	42
5.6.4 Example 4	42
5.7 API	43
5.7.1 Overview	43
5.7.2 Configuration/Clear API	44
5.7.2.1 configure-bst-feature	44
5.7.2.2 configure-bst-tracking	45
5.7.2.3 configure-bst-thresholds	47
5.7.2.4 configure-bst-multi-thresholds	48
5.7.2.5 clear-bst-statistics	49
5.7.2.6 clear-bst-thresholds	50
5.7.2.7 clear-bst-congestion-drop-counters	50
5.7.3 Status/Reporting API	51
5.7.3.1 get-bst-feature	51
5.7.3.2 get-bst-tracking	52
5.7.3.3 get-bst-thresholds	53
5.7.3.4 get-bst-report	55
5.7.3.5 get-bst-congestion-drop-counters	57
5.7.4 Trigger Reports	63
Chapter 6: Packet Trace	65
6.1 Overview	65
6.2 Configurability	65
6.3 Packet Format	66
6.4 Live Traffic Triggered Trace-Profile/Drop-Profile	67
6.5 Timestamps	68
6.6 Deciphering Trace-Profile	68

6.7 Deciphering Drop-Profile	70
6.8 API	71
6.8.1 Overview	71
6.8.2 Configuration/Clear API	71
6.8.2.1 configure-packet-trace-feature	71
6.8.3 Status/Reporting API	72
6.8.3.1 get-packet-trace-feature	72
6.8.3.2 get-packet-trace-lag-resolution	73
6.8.3.3 get-packet-trace-ecmp-resolution	77
6.8.3.4 get-packet-trace-profile	80
6.8.3.5 get-packet-drop-reason	84
6.8.3.6 get-packet-dropctrs	85
Chapter 7: Black Hole Detection	87
7.1 Overview	87
7.2 Black Hole Ports	87
7.3 Reporting Black Hole Events	88
7.4 Configurability	88
7.5 API	89
7.5.1 Overview	89
7.5.2 Configuration API	89
7.5.2.1 black-hole-detection-enable	89
7.5.2.2 configure-black-hole	90
7.5.2.3 cancel-black-hole	92
7.5.3 Status/Reporting API	92
7.5.3.1 get-black-hole-detection-enable	92
7.5.3.2 get-black-hole	93
7.5.3.3 get-sflow-sampling-status	94
7.5.4 Notifications	95
7.5.4.1 get-black-hole-event-report	95
Chapter 8: Flow Tracker	96
8.1 Overview	96
8.2 Communication Exchanges	96
8.2.1 Information	96
8.2.2 Entities	97
8.3 IPFIX Templates	98
8.3.1 Flow-Group Record Data Set Template	98
8.3.2 Flow Record Data Set Template	99
8.4 Configurability	99
8.5 API	100

8.5.1 Overview	100
8.5.2 Configuration API	101
8.5.2.1 configure-ft-feature	101
8.5.2.2 create-ft-collector	102
8.5.2.3 remove-ft-collector	104
8.5.2.4 create-ft-flowgroup	104
8.5.2.5 remove-ft-flowgroup	109
8.5.2.6 clear-ft-flowgroup-statistics	109
8.5.3 Status/Reporting API	110
8.5.3.1 get-ft-collector	110
8.5.3.2 get-ft-flowgroup	111
8.5.3.3 get-ft-flowgroup-statistics	112
8.5.3.4 get-ft-capabilities	114
8.5.3.5 get-ft-status	115
8.6 Configuration Persistence	116
Chapter 9: Reference Applications	117
Chapter 10: License	118
Chapter 11: Source Code	119
Chapter 12: Documentation	120
Appendix A: JSON Payload Revision History.....	121
Revision History	124
BroadView-SP101; May 31, 2018	124
BroadView-SP100; April 26, 2017	124

Chapter 1: Introduction

1.1 Overview

With the proliferation of Software-Defined Networking (SDN), multi-tenant networks, and server virtualization—aided by cloud deployments for applications and storage—network complexity is growing exponentially. Troubleshooting such networks has become an increasingly daunting task. Network operators need increased visibility into the network and deeper telemetry data in order to remain in control of the network and to ensure optimal network resource utilization. BroadView™ Instrumentation software provides this critical visibility and telemetry information.

Broadview Instrumentation Agent unlocks the potential of Broadcom® switch silicon by augmenting CPU functionality to deliver advanced network analytics. The Agent source code is Apache licensed, sports a modular design, offers REST API for interaction, and is easily adaptable into customer solutions.

The Broadview Instrumentation Agent communicates with the underlying Broadcom switch silicon. It collects various telemetry and visibility information, runs algorithms on the data, packages the information appropriately, and provides it to registered clients. Similarly, the Agent configures the silicon based on the configuration requests from the client. On select devices, BroadView Instrumentation Agent supports sending sampled packets to an sFlow receiver on the network¹

The Agent communicates with clients using REST-style communication, with the data exchange in JSON-RPC (2.0) format. The Agent supports both the pull model of operation (the client requests data and obtains it) as well as the push model of operation (the Agent sends periodic reports, asynchronously).

The Agent works with a wide range of Broadcom switch silicon, and the features offered by the Agent operate in accordance with the underlying silicon.

In a deployment scenario, the BroadView Agent works in conjunction with the Network Operating System (NOS). The Agent can work completely integrated into the NOS, or as an independent application on the switch. Likewise, the telemetry information can be used by the NOS itself and/or by clients interacting with the Agent via the REST API.

This document describes the Agent architecture and each of the supported BroadView features. Each BroadView feature description is followed by a listing of REST APIs supported by the feature.

1. It may be noted that the actual sFlow functionality is provided by the silicon. The Agent enables this functionality upon a Client request. Consult the user documentation of the silicon for the details such as how sample packets are packaged into sFlow frames, packet formats etc. Such details are beyond the scope of this document.

1.2 Supported Features by Silicon²

Table 1 lists the Broadview Instrumentation Agent features and identifies which Broadcom devices support each feature.

Table 1: Supported Features by Silicon

Feature	Supported Silicon
Buffer Statistics and Tracking (BST)	BCM56850 BCM56760 BCM56860 BCM56870 BCM56960 BCM56965
Packet Trace (PT)	BCM56870 BCM56960 BCM56965
Black Hole Detection (BHD) ^a	BCM56850 BCM56960 BCM56965 BCM88375
Flow Tracker (FT)	BCM56960 BCM56870

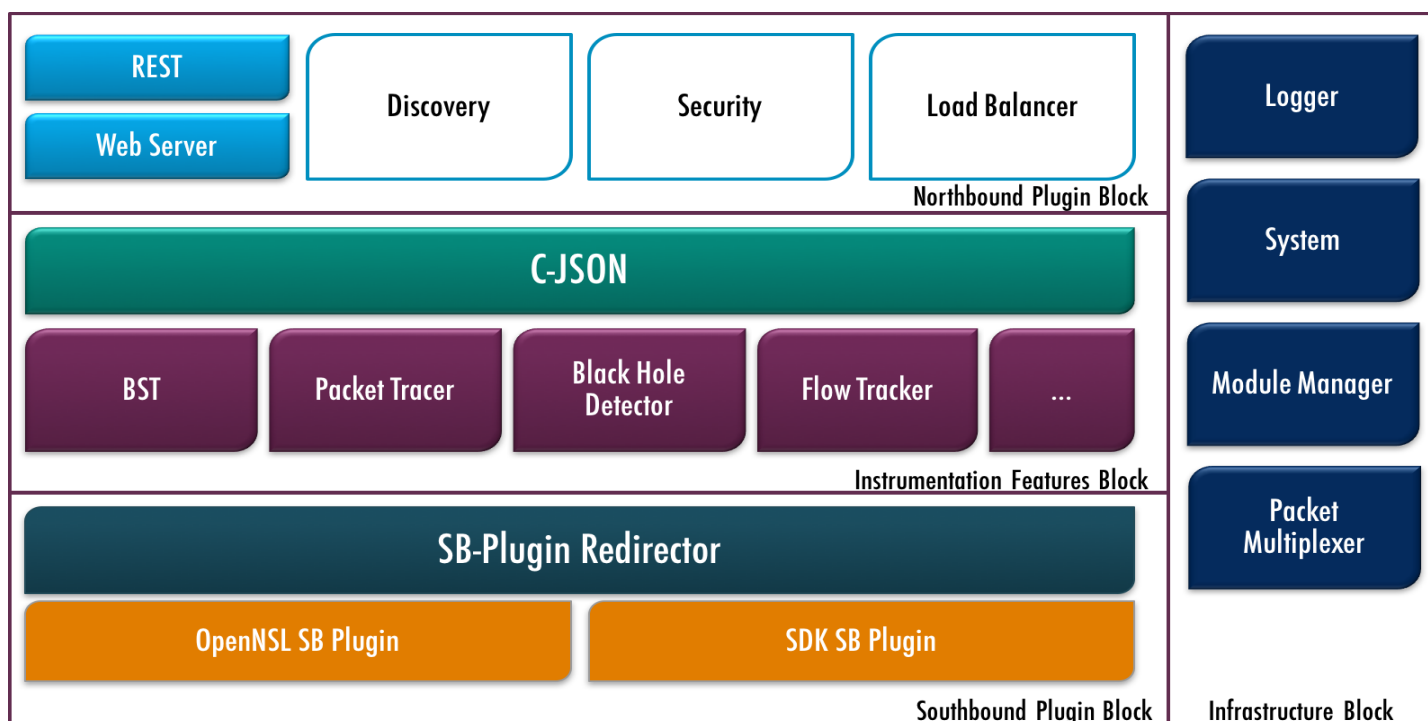
a. BHD sFlow sampling is supported only on the BCM56960.

- It is possible that not all features and silicon devices are supported in a given release of the Broadview Instrumentation Agent. For details, consult the release collateral.

Chapter 2: Architecture

The Broadview Instrumentation Agent has a modular architecture and consists of a set of components that serve different purposes in the architecture. Some core components implement the telemetry and analytics algorithms, while others function as infrastructure components. Some of the components are meant for porting and/or customization, and the reference implementations of those components are for illustrative purposes only. For example, by default, the Agent software does not support any authentication or security for the communication with the collector. However, the security and web server modules can be easily enhanced or added to provide additional functionality. Subsequent pages list various customizations that are possible.

Figure 1: BroadView Agent Architecture



2.1 Architecture Blocks

The components that form the Agent software are divided into the following four blocks described below:

- [Northbound Plug-in Block](#)
- [Infrastructure Block](#)
- [Southbound Plug-in Block](#)
- [Instrumentation Features Block](#)

2.1.1 Northbound Plug-in Block

The Northbound (NB) Plug-in block is responsible for setting up and managing the communication with the applications/controller(s)/collector(s), referred to as collectors, that interface with the Agent. The distribution includes a reference application that exercises the NB API of the reference Agent.

The Agent uses REST-style communication to communicate with the collector, with the data exchange in the JSON-RPC format. The REST and web server components handle these interactions with the collectors. The REST component extracts the incoming JSON-RPC request, deciphers the API being invoked by the collector, and passes on the JSON message to the appropriate Instrumentation Features block component (using the Module Manager component services).

The Discovery component is responsible for detecting the presence of one or more collectors on the network and establishing communication between the Agent and the device. It also maintains a heartbeat channel with the collector for detecting any loss of communication.

The Security component provides encryption support for the communication between the Agent and the collector. It also provides the necessary authentication for the discovered collectors.

The Load Balancer component provides active and standby collectors. It keeps track of communication errors with the designated active collector. In the event of persistent failures, this component switches the standby collector to the active state and establishes communication with it.

2.1.1.1 Reference Implementation

The reference implementation of the Discovery module reads the collector IP addresses and communication port numbers from a predefined configuration file.

The Security module is not provided as part of the reference implementation. This means that the Agent allows unencrypted communication between the collector and the Agent. Likewise, the collectors are considered genuine, and no attempt is made to validate the authenticity of the collector.

The Load Balancer module is not provided as part of the reference implementation. Any errors during either the pull or during the push operations are logged, and communication is reattempted.

2.1.2 Infrastructure Block

The Infrastructure block provides the necessary infrastructure and utilities for the other modules in the Agent.

The Logging component provides logging functionality to the Agent system. It provides API functions that can be used for logging various events and data. Different levels of logging are supported. The Module Manager component provides for registration of various Instrumentation components with the Agent. Each of the Instrumentation components must register with the Module Manager to provide the details of various APIs it supports, the feature identification, supported silicon, and the callback functions that process the API. When a REST API is invoked by the collector, the Module Manager component assists the REST component in invoking the appropriate handler for the API, based on the registration information.

The Instrumentation components are typically C code linked with the Agent. The components register the supported REST API with the Module Manager at run time during the initialization process.

The System component is in charge of setting up the Agent based on the configuration and starting/stopping various other components as needed. The System component also provides the timer API, which can be used by the other components for periodic data collection by the hardware. Note that the periodic data collection may be disabled/enabled by the collector on a system-wide basis or selectively on a per-component basis.

2.1.2.1 Reference Implementation

The reference implementation of the Logging component uses syslog. It can easily be enhanced to use a simple SQLite3 database.

2.1.3 Southbound Plug-in Block

The Southbound block features different plug-in modules that conform to the SB-API specification. Each of the plug-ins registers with the SB-Plug-in Redirector component, with the list of features as well as the silicon that are supported by the plug-in. When any BroadView Instrumentation component attempts to communicate with the silicon, it invokes the corresponding API of the SB-Plug-in Redirector component, which in turn uses the registration information to invoke the corresponding plug-in component. The purpose of the SB-Plug-in is to allow functions to be written that are specific to the mechanism that is available to obtain the instrumentation information from the silicon on any specific system. For example, System 1 could allow the use of an SDK API while System 2 may have its own API (such as the OpenNSL API).

The SB-Plug-in modules are typically shared object libraries linked with the Agent. The plug-in registration is done by invoking the SB-Plug-in Redirector component's registration API.

2.1.3.1 Reference Implementation

The reference implementation consists of an SDK SB Plug-in and OpenNSL SB Plug-in.

2.1.4 Instrumentation Features Block

The Instrumentation components are registered statically with the Agent. Upon startup, the Agent initializes each of the registered components as part of the global initialization sequence. The components are required to register the supported silicon, features, and the REST API with the Module Manager component. Such registration enables the Module Manager to understand the set of REST APIs that are supported by the entire Agent. Any feature configuration REST API is also included in the registration.

The web server thread invokes the API handler function with the JSON message buffer.

The supported data formats for the API (commands) as well as the responses are documented in XML format. A tool is run on these XML files, which can generate C code that can parse the JSON buffer and convert the incoming data to an appropriate C structure. This generated code is used as the implementation for the handler functions. After the C structure is derived, the control is handed over to the component thread. The web server thread acknowledges the REST API to the collector.

The component, under its own thread context now, makes the appropriate function calls to process the request and invoke the appropriate SB-Plug-in API calls. It is also likely that the component may not be required to make any SB-Plug-in calls for fulfilling the REST API since the data may be collected by the periodic thread already. When all the required data is available, the data is converted into JSON. Using a web server API function, the data is sent back to the collector.

2.2 Execution in Linux

Various components of the Agent are compiled and linked into a single Linux-executable³. This user-space executable (as a single process) runs multiple threads within, based on the needs.

Typically, each of the Instrumentation feature modules start and maintain a thread. The API invoked by the collector is to be executed in the corresponding feature/module thread context. In addition, there is a system-wide timer thread meant for the periodic data collection. The other modules register their data-gathering functions with the system data-collection thread and suggest the periodicity with which the functions need to be invoked. The registered callbacks are invoked in the data-collection thread context, and care must be taken by the callbacks to ensure data integrity.

The components in the Agent use POSIX API for any required inter-thread communication.

No kernel space operations are performed within the Agent process. Multiple Agent instances running in a single Linux environment is not supported.

3. Linux Operating System is assumed.

2.3 Using the Agent with a NOS

This section describes high-level porting considerations when adapting the Agent to a NOS.

2.3.1 SB Plug-ins

A Southbound plug-in provides a means of interaction with the underlying silicon to the rest of the Agent software. The reference Agent has an SB plug-in that uses the SDK /OpenNSL API calls directly inside the implementation. For various reasons, this may not be desirable when the Agent is used in conjunction with a NOS. It may be preferable to route the silicon accesses via the NOS.

If the access to the silicon is to be routed via the NOS, an appropriate SB plug-in must be developed.

2.3.2 NB Plug-ins

The default web server provided by the Agent does not handle the security or the discovery/load-sharing aspects. If this is not an ideal implementation, use one of the following alternatives:

- Replace the default web server with another web server component or integrate the Agent into the NOS-supported web server.
- Use the NOS web server to handle the incoming REST requests. After authenticating the request, the NOS web server hands the request over to the Agent web server.

Chapter 3: Communication

3.1 REST

The BroadView Agent supports a REST-based API and uses JSON-RPC messaging as the payload over REST. The name of the API becomes the method in the JSON-RPC message, and any associated parameters for the command form the **params** in the JSON-RPC message.

All the commands are to be sent as HTTP 1.1 POST operations, with the JSON content as the HTTP message body.

By default, the web server provided as part of the reference Agent uses the following URL scheme to provide the access to the API:

<http://<ipaddr>:8080/broadview/<feature>/<api-name>>

For example, the `configure-bst-feature` API is accessible via the following URL:

<http://<ipaddr>:8080/broadview/bst/configure-bst-feature>

For the Generic Agent Management API, the URL scheme does not have the feature name embedded in the URL. For example, the `configure-system-feature` API is accessible via the following URL:

<http://<ipaddr>:8080/broadview/configure-system-feature>

3.2 JSON-RPC Messaging

3.2.1 Request Response Messages

The API supported by BroadView Agent uses JSON-RPC messaging as payload over HTTP. There are three forms of JSON-RPC messages that are supported by the Agent:

- **Request** (from the collector and received/processed by the Agent), sometimes also called as a **Command**.
- **Response** (from the Agent and received/processed by the collector in response to a previous request).
- **Notification** (from the Agent and received/processed by the collector, asynchronously).

An example JSON-RPC **request** message sample is provided below for illustrative purposes.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-feature",
  "params": {},
  "asic-id": "1",
  "id": 1232
}
```

A sample JSON-RPC **response** message is provided below for illustrative purposes.

```
{
  "jsonrpc": "2.0",
  "asic-id": "1",
  "version": "1",
  "result": {
    "bst-enable": 1,
    "bst-trackers": {
      "device" : 0,
      "egress-unicast-queue" : 1
    }
  },
  "id": 1232
}
```

A notification message is similar to a request message, with the exception that it does not have an “id” field. More details are available at <http://www.jsonrpc.org/specification>.

When an Agent receives a request from the collector, the **method** field specifies the request that is sent. Any data associated with the request is part of the **params** field.

3.2.2 Time Stamps

The messages sent by the Agent may contain the timestamp indicating the time at which the data is read from the ASIC. The timestamp is reported in the YYYY-MM-DD - HH:MM:SS format.

A sample JSON report snippet containing a timestamp is provided below.

```
...
"method": "get-bst-report",
"asic-id": "2",
"time-stamp": "2014-10-18 - 00:15:04 ",
"report": [{
  ...
```

3.2.3 Version

The response message from the Agent contains a field called *version*. This field contains an integer indicating the specific JSON payload version being used by the Agent. The version number is incremented for every set of changes introduced in the JSON payload. The client may use this version number information to parse/decipher the response messages from the Agent.

[Appendix A, JSON Payload Revision History](#) lists the changes introduced in the JSON payload with various versions.

3.3 Multi-ASIC Platforms

The Agent supports platforms with more than one switch ASIC device. The configuration command/reporting is on a per-ASIC basis. The configuration commands, responses, and notification reports carry the ASIC identifier as part of the message.

All parts of the REST API supported by the Broadview Instrumentation Agent take `asic-id` as a parameter via the JSON requests. Similarly, all the responses and notification messages from the Agent include the `asic-id` as a parameter in the response JSON. This is illustrated below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-feature",
  "params": { },
  "asic-id": "1",
  "id": 1232
}

...
"method": "get-bst-report",
"asic-id": "2",
"time-stamp": "2014-10-18 - 00:15:04 ",
"report": [{
  ...
```

For a single ASIC platform, the ASIC identifier (`asic-id`) part of the configuration command for response messages and notification reports is ignored by the Agent, and the corresponding value is set to 0.

This document describes the REST API for all supported features in detail, including specific parameters for the request/response messages. These descriptions do not include the `asic-id` parameter to avoid repetitive duplication.

3.4 Error Reporting

The Agent uses the standard HTTP error reporting mechanism for reporting any errors. The error codes are listed in [Table 2](#).

Table 2: Error Codes

HTTP Status Code	Description
HTTP 200	Indicates that the request successfully executed on the Agent. If the request is for status/data, the response accompanies this status. ^a
HTTP 204	Indicates that the Agent has successfully executed the request, but no content is being returned.
HTTP 400	Indicates a JSON error in the request message. The JSON message is badly formatted and contains errors.
HTTP 403	Indicates an error. The method cannot be invoked.
HTTP 404	Indicates an error. The method requested is not supported.
HTTP 500	Indicates an error. There was an internal error on the Agent that caused this.

- a. An HTTP 200 status indicates successful execution only at the transport and agent level, but not necessarily at the application/feature level. In most cases, they are the same. However, in some corner cases, they may be different. Where applicable, the JSON response messages indicate application-level execution status, and the client is required to parse the response.

Each of the errors listed above (HTTP error codes 4xx and 5xx) have a JSONRPC payload in the response packet to give additional error details in the JSON payload. An example of an HTTP 404 error is provided below.

HTTP/1.1 404 Not Found

Server: BroadViewAgent (Unix) (Linux)

Content-Type: text/json

```
{
  "jsonrpc": "2.0",
  "version": "1",
  "error": {
    "code": -32601,
    "message": "Method not found"
  },
  "id": 1
}
```

Chapter 4: Generic Agent Management

4.1 Overview

The Broadview Instrumentation Agent offers an API that helps clients and SDN controllers discover switches running the Agent and determine the switch capabilities, including both ASIC capabilities and the features supported by the Agent.

4.2 Heartbeat Messaging

Heartbeat messages are keep-alive messages sent by the Agent to the client. The heartbeat messaging allows clients to auto detect the switches running the Agent and to configure the Agent, without requiring manual intervention. The heartbeat message is in the form of REST/JSON notification messages.

Upon configuration, the Agent sends periodic heartbeat messages to a client. The messages contain the Agent and ASIC capabilities. By default, the heartbeat⁴ messages are sent every five seconds. Upon discovering the Agent, the client can start configuring the Agent and use the BroadView capabilities for monitoring the switch/network. Optionally, the Agent can reduce the frequency with which the heartbeat messages are sent to the client.

The heartbeat mechanism allows the registration of a specific switch with BroadView Agent capability but does not provide information about any connectivity of the switch to other switches.

4. Sometimes the heartbeat messages prior to Client configuration of the Agent are called registration/discovery messages. The content of the message is the same, independent of the name.

4.3 Switch Properties

The switch properties are a set of parameters describing the Agent capabilities as well as the switch information. This set of parameters packed inside a JSON message becomes the heartbeat message. It can also be retrieved on-demand via the `get-switch-properties` API.

The parameters are described in [Table 3](#).

Table 3: Switch Properties Parameters

Parameter	Type	Description
number-of-asics	Integer	Number of ASICs on the switch.
asic-info	Multi-parameter	Describes each of the ASICs onboard. More details are below in Table 4 .
supported-features	Array of strings	A list of strings indicating the features supported by the Agent.
network-os	String	The NOS currently used on the switch.
uid	String	Unique identifier for this switch. This unique ID is the key for the SDN controller to map the switch to the nodes existing in their discovery database.
agent-ip	String	IP address of the switch where the Agent is being run.
agent-port	String	TCP port number of the switch, at which the Agent is listening.
agent-sw-version	String	Software version number for the Agent.
last-saved-configuration	String	Timestamp indicating the time the configuration was last saved by user.
agent-up-time	Integer	Number of seconds since Agent has started functioning.
unsaved-config	Boolean	Indicates whether user has changed the configuration of the agent, but has not saved it to persistent storage.

The `asic-info` parameter is a multi-parameter, and each of the sub-parameters is described in [Table 4](#). The `asic-info` parameter is encoded as a positional-parameter list in the JSON message.

Table 4: asic-info Subparameters

Position	Parameter	Type	Description
1	asic-id	String	Identifier for an ASIC on the switch. This ID must be used in all subsequent interactions with the Agent to refer to this ASIC.
2	chip-id	String	Indicates the part number of the silicon.
3	num-ports	Integer	Number of ports available on the switch, managed by this ASIC.

A sample heartbeat message is provided below.

```
{
  "jsonrpc": "2.0",
  "method": "get-switch-properties",
  "version": "2",
  "time-stamp": "2015-10-18 - 00:15:04",
  "result": {
    "number-of-asics": 1,
    "asic-info": [
      [
        "1",
        "BCM56850",
        78
      ]
    ],
    "supported-features": [
      "BST"
    ],
    "network-os": "openNSL",
    "uid": "0000d80bb99bbbbbb",
    "agent-ip": "192.168.1.2",
    "agent-port": "8080",
    "agent-sw-version": "3.0.0.1",
    "last-saved-configuration": "2017-01-01 - 11:43:07",
    "agent-up-time": 52,
    "unsaved-config" : 1
  },
  "id": 10
}
```

4.4 Configurability

Configurability is provided under the Generic Agent Management category.

The Agent can be configured to enable/disable heartbeat messaging to the client, as well as the interval at which the registration messages are to be sent.

The Agent can be asked store the current configuration in persistent storage. Agent can also be asked to clear/erase the current configuration and restore the factory defaults, or restore to the previously saved configuration. The Agent uses NOS services for storing/retrieving the configuration blob to/from a persistent storage media. The reference implementation stores the configuration blob to a local file. During the Agent startup sequence, if a valid persistent configuration is found, that configuration is automatically applied.

4.5 API

4.5.1 Overview

This section lists various APIs supported for the Generic Agent Management. A description of each API is in subsequent sections.

Table 5: APIs for Generic Agent Management

API	Type	Description
configure-system-feature	Configuration	Configure global Agent management functionality.
save-configuration	Configuration	Stores the configuration persistently.
restore-configuration	Configuration	Clears the current configuration and restores either the factory defaults or the previously saved configuration.
get-configuration	Status/Reporting	Returns the stored configuration
get-port-queue-map	Status/Reporting	Obtain the hardware-queue to user-queue mapping.
get-system-feature	Status/Reporting	Retrieve Agent management configuration.
get-switch-properties	Status/Reporting	Retrieve the switch capabilities.
cancel-request^a	Configuration	Stop a previously initiated periodic reporting command.

a. This API is not specific to any single feature, and may be used to cancel any previously initiated periodic reporting command.

4.5.2 Configuration API

4.5.2.1 configure-system-feature

The `configure-system-feature` API sets up the core Agent functionality on the switch. The parameters associated with this command are described in [Table 6](#).

Table 6: configure-system-feature Parameters

Parameter	Type	Description
<code>heartbeat-enable</code>	Boolean	When enabled, the Agent asynchronously sends the registration and heartbeat message to the collector. By default, the sending discovery message is turned on.
<code>msg-interval</code>	Integer (range 1 to 600)	Determines the interval in seconds with which the registration and heartbeat messages are sent to the collector. The default value is 5 seconds.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "configure-system-feature",
  "params": {
    "heartbeat-enable": 1,
    "msg-interval": 10
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

4.5.2.2 save-configuration

The `save-configuration` API stores the Agent's current configuration in a persistent medium. Any existing persistent configuration is overwritten. There are no parameters associated with this command.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "save-configuration",
  "params": {},
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

4.5.2.3 restore-configuration

The `restore-configuration` API clears the Agent's current configuration and restores either the factory defaults or the configuration from the persistent medium. The parameters associated with this command are described below.

Table 7: restore-configuration API Parameter

Parameter	Type	Description
<code>type</code>	String Enumeration	Indicates the configuration to be used after clearing the current configuration. Supported values are <i>factory-reset</i> and <i>saved-configuration</i> .

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "restore-configuration",
  "params": {
    "type": "factory-reset"
  },
  "id": 2
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

4.5.2.4 cancel-request

The `cancel-request` API is a generic API, and it cancels a previously configured periodic reporting request, such as `get-bst-congestion-drop-counters` on the Agent. The parameter associated with this command is described in [Table 8](#).

Table 8: cancel-request API Parameter

Parameter	Type	Description
<code>request-id</code>	Integer	The <code>request-id</code> parameter is the ID of the request made already by the client. If this parameter value is zero (0), then the agent cancels all previously initiated periodic reporting requests.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "cancel-request",
  "asic-id": "1",
  "params": {
    "request-id": 7
  },
  "id": 4
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

4.5.3 Status/Reporting API

4.5.3.1 get-configuration

The `get-configuration` API retrieves the Agent's specified configuration. The parameters associated with this command are described below.

Table 9: restore-configuration API Parameters

Parameter	Type	Description
<code>type</code>	String Enumeration	Indicates the configuration to be retrieved. Only supported value is <i>saved</i> , this returns the configuration saved in persistent medium.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-configuration",
  "params": {
    "type": "saved"
  },
  "id": 2
}
```

The Agent returns the configuration or returns an appropriate error code if needed.

The returned configuration is a JSON formatted string. The configuration is made-up of individual JSON commands (sent by user via REST API) concatenated into a JSON array. A sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-configuration",
  "version": "6",
  "commands": [
    {
      "jsonrpc": "2.0",
      "method": "configure-bst-feature",
      "asic-id": "1",
      "params": {
        "bst-enable": 0,
        "collection-interval": 60,
        "send-async-reports": 0,
        "stats-in-percentage": 0,
        "stat-units-in-cells": 0,
        "trigger-rate-limit": 1,
        "send-snapshot-on-trigger": 1,
        "trigger-rate-limit-interval": 1,
        "async-full-reports": 0
      },
      "id": 1
    },
    {
      "jsonrpc": "2.0",
      "method": "configure-bst-tracking",
      "asic-id": "1",

```



```

    "params": {
      "track-peak-stats": 0,
      "track-ingress-port-priority-group": 1,
      "track-ingress-port-service-pool": 1,
      "track-ingress-service-pool": 1,
      "track-egress-port-service-pool": 1,
      "track-egress-service-pool": 1,
      "track-egress-uc-queue": 1,
      "track-egress-uc-queue-group": 1,
      "track-egress-mc-queue": 1,
      "track-egress-cpu-queue": 1,
      "track-egress-rqe-queue": 1,
      "track-device": 1
    },
    "id": 1
  },
  {
    "jsonrpc": "2.0",
    "method": "configure-bst-multi-thresholds",
    "asic-id": "1",
    "params": [
      {
        "realm": "device",
        "data": 1000
      },
      {
        "realm": "ingress-port-priority-group",
        "data": []
      },
      {
        "realm": "ingress-port-service-pool",
        "data": []
      },
      {
        "realm": "ingress-service-pool",
        "data": []
      },
      {
        "realm": "egress-cpu-queue",
        "data": []
      },
      {
        "realm": "egress-mc-queue",
        "data": []
      },
      {
        "realm": "egress-port-service-pool",
        "data": []
      },
      {
        "realm": "egress-rqe-queue",
        "data": []
      },
      {
        "realm": "egress-service-pool",
        "data": []
      }
    ]
  }

```

```

        "realm": "egress-uc-queue",
        "data": []
      },
      {
        "realm": "egress-uc-queue-group",
        "data": []
      }
    ],
    "id": 1
  },
  {
    "jsonrpc": "2.0",
    "method": "configure-packet-trace-feature",
    "asic-id": "1",
    "params": {
      "packet-trace-enable": 0
    },
    "id": 1
  },
  {
    "jsonrpc": "2.0",
    "method": "configure-system-feature",
    "params": {
      "heartbeat-enable": 1,
      "msg-interval": 5
    },
    "id": 1
  }
],
"id": 2
}

```

4.5.3.2 get-port-queue-map

The `get-port-queue-map` command is used to retrieve the hardware queue to user-assigned queue mapping for every port. There are no parameters for this command.

A sample JSON-RPC request to get the port to queue mapping is shown below.

```

{
  "jsonrpc": "2.0",
  "method": "get-port-queue-map",
  "asic-id": "1",
  "params": { },
  "id": 1
}

```

The response parameters are as shown in [Table 10](#). These parameters are provided for both unicast queues and multicast queues for each of the ports.

Table 10: get-port-queue-map Response Parameters

Parameter	Type	Description
hw-queue	Integer	The queue number as supported by the hardware
user-queue	Integer	The queue number for the specified port, for which the hw-queue is assigned to.

A sample port-queue map for the above request is as follows:

```
{
  "jsonrpc": "2.0",
  "method": "get-port-queue-map",
  "asic-id": "1",
  "version": "4",
  "result": [
    {
      "port": "1",
      "unicast-queue-map": [
        {
          "hw-queue": 0,
          "user-queue": 0
        },
        {
          "hw-queue": 1,
          "user-queue": 1
        },
        {
          "hw-queue": 2,
          "user-queue": 2
        },
        {
          "hw-queue": 3,
          "user-queue": 3
        },
        {
          "hw-queue": 4,
          "user-queue": 4
        },
        {
          "hw-queue": 5,
          "user-queue": 5
        },
        {
          "hw-queue": 6,
          "user-queue": 6
        },
        {
          "hw-queue": 7,
          "user-queue": 7
        }
      ],
      "multicast-queue-map": [
        {
          "hw-queue": 0,
          "user-queue": 0
        },
        {
          "hw-queue": 1,
          "user-queue": 1
        },
        {
          "hw-queue": 2,
          "user-queue": 2
        },
        {
          "hw-queue": 3,
```

```
        "user-queue": 3
      },
      {
        "hw-queue": 4,
        "user-queue": 4
      },
      {
        "hw-queue": 5,
        "user-queue": 5
      },
      {
        "hw-queue": 6,
        "user-queue": 6
      },
      {
        "hw-queue": 7,
        "user-queue": 7
      }
    ]
  },
  {
    "port": "2",
    "unicast-queue-map": [
      {
        "hw-queue": 8,
        "user-queue": 0
      },
      {
        "hw-queue": 9,
        "user-queue": 1
      },
      {
        "hw-queue": 10,
        "user-queue": 2
      },
      {
        "hw-queue": 11,
        "user-queue": 3
      },
      {
        "hw-queue": 12,
        "user-queue": 4
      },
      {
        "hw-queue": 13,
        "user-queue": 5
      },
      {
        "hw-queue": 14,
        "user-queue": 6
      },
      {
        "hw-queue": 15,
        "user-queue": 7
      }
    ],
    "multicast-queue-map": [
      {
```

```
        "hw-queue": 8,
        "user-queue": 0
    },
    {
        "hw-queue": 9,
        "user-queue": 1
    },
    {
        "hw-queue": 10,
        "user-queue": 2
    },
    {
        "hw-queue": 11,
        "user-queue": 3
    },
    {
        "hw-queue": 12,
        "user-queue": 4
    },
    {
        "hw-queue": 13,
        "user-queue": 5
    },
    {
        "hw-queue": 14,
        "user-queue": 6
    },
    {
        "hw-queue": 15,
        "user-queue": 7
    }
]
},
{
    "port": "3",
    "unicast-queue-map": [
        {
            "hw-queue": 16,
            "user-queue": 0
        },
        {
            "hw-queue": 17,
            "user-queue": 1
        },
        {
            "hw-queue": 18,
            "user-queue": 2
        },
        {
            "hw-queue": 19,
            "user-queue": 3
        },
        {
            "hw-queue": 20,
            "user-queue": 4
        },
        {
            "hw-queue": 21,
```

```
        "user-queue": 5
      },
      {
        "hw-queue": 22,
        "user-queue": 6
      },
      {
        "hw-queue": 23,
        "user-queue": 7
      }
    ],
    "multicast-queue-map": [
      {
        "hw-queue": 16,
        "user-queue": 0
      },
      {
        "hw-queue": 17,
        "user-queue": 1
      },
      {
        "hw-queue": 18,
        "user-queue": 2
      },
      {
        "hw-queue": 19,
        "user-queue": 3
      },
      {
        "hw-queue": 20,
        "user-queue": 4
      },
      {
        "hw-queue": 21,
        "user-queue": 5
      },
      {
        "hw-queue": 22,
        "user-queue": 6
      },
      {
        "hw-queue": 23,
        "user-queue": 7
      }
    ]
  },
  {
    "port": "4",
    "unicast-queue-map": [
      {
        "hw-queue": 24,
        "user-queue": 0
      },
      {
        "hw-queue": 25,
        "user-queue": 1
      },
      {

```

```
        "hw-queue": 26,
        "user-queue": 2
    },
    {
        "hw-queue": 27,
        "user-queue": 3
    },
    {
        "hw-queue": 28,
        "user-queue": 4
    },
    {
        "hw-queue": 29,
        "user-queue": 5
    },
    {
        "hw-queue": 30,
        "user-queue": 6
    },
    {
        "hw-queue": 31,
        "user-queue": 7
    }
],
"multicast-queue-map": [
    {
        "hw-queue": 24,
        "user-queue": 0
    },
    {
        "hw-queue": 25,
        "user-queue": 1
    },
    {
        "hw-queue": 26,
        "user-queue": 2
    },
    {
        "hw-queue": 27,
        "user-queue": 3
    },
    {
        "hw-queue": 28,
        "user-queue": 4
    },
    {
        "hw-queue": 29,
        "user-queue": 5
    },
    {
        "hw-queue": 30,
        "user-queue": 6
    },
    {
        "hw-queue": 31,
        "user-queue": 7
    }
]
```

```
},
{
  "port": "5",
  "unicast-queue-map": [
    {
      "hw-queue": 32,
      "user-queue": 0
    },
    {
      "hw-queue": 33,
      "user-queue": 1
    },
    {
      "hw-queue": 34,
      "user-queue": 2
    },
    {
      "hw-queue": 35,
      "user-queue": 3
    },
    {
      "hw-queue": 36,
      "user-queue": 4
    },
    {
      "hw-queue": 37,
      "user-queue": 5
    },
    {
      "hw-queue": 38,
      "user-queue": 6
    },
    {
      "hw-queue": 39,
      "user-queue": 7
    }
  ],
  "multicast-queue-map": [
    {
      "hw-queue": 32,
      "user-queue": 0
    },
    {
      "hw-queue": 33,
      "user-queue": 1
    },
    {
      "hw-queue": 34,
      "user-queue": 2
    },
    {
      "hw-queue": 35,
      "user-queue": 3
    },
    {
      "hw-queue": 36,
      "user-queue": 4
    },
  ],
```



```
{
  {
    "hw-queue": 37,
    "user-queue": 5
  },
  {
    "hw-queue": 38,
    "user-queue": 6
  },
  {
    "hw-queue": 39,
    "user-queue": 7
  }
]
},
{
  "port": "6",
  "unicast-queue-map": [
    {
      "hw-queue": 40,
      "user-queue": 0
    },
    {
      "hw-queue": 41,
      "user-queue": 1
    },
    {
      "hw-queue": 42,
      "user-queue": 2
    },
    {
      "hw-queue": 43,
      "user-queue": 3
    },
    {
      "hw-queue": 44,
      "user-queue": 4
    },
    {
      "hw-queue": 45,
      "user-queue": 5
    },
    {
      "hw-queue": 46,
      "user-queue": 6
    },
    {
      "hw-queue": 47,
      "user-queue": 7
    }
  ],
  "multicast-queue-map": [
    {
      "hw-queue": 40,
      "user-queue": 0
    },
    {
      "hw-queue": 41,
      "user-queue": 1
    }
  ]
}
```

```
    },
    {
      "hw-queue": 42,
      "user-queue": 2
    },
    {
      "hw-queue": 43,
      "user-queue": 3
    },
    {
      "hw-queue": 44,
      "user-queue": 4
    },
    {
      "hw-queue": 45,
      "user-queue": 5
    },
    {
      "hw-queue": 46,
      "user-queue": 6
    },
    {
      "hw-queue": 47,
      "user-queue": 7
    }
  ]
},
" id": 1
}
```

The Agent returns the requested map immediately.

The Agent acknowledges the command and returns an appropriate error code if needed.

4.5.3.3 get-system-feature

The `get-system-feature` API retrieves the current Agent configuration. There is no parameter associated with this API. The parameters that form the response message are the same as described in the [configure-system-feature](#) API.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-system-feature",
  "params": { },
  "id": 1
}
```

An associated sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-system-feature",
  "version": "2",
  "result": {
    "heartbeat-enable": 1,
    "msg-interval": 5
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

4.5.3.4 get-switch-properties

The `get-switch-properties` API retrieves the switch properties as described in [Section 4.3, Switch Properties](#). There are no parameters associated with this command.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-switch-properties",
  "asic-id": "1",
  "params": { },
  "id": 1
}
```

An associated sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-switch-properties",
  "asic-id": "1",
  "params": { },
  "id": 1
}
```

An associated sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-switch-properties",
  "version": "6",
  "time-stamp": "2017-10-18 - 08:15:04",
  "result": {
    "number-of-asics": 1,
    "asic-info": [
      [
        "1",
        "BCM56850",
        78
      ]
    ],
    "supported-features": [
      "BST", "PT", "PDM"
    ],
    "network-os": "SDK",
    "uid": "0000d80bb99bbbbbb",
    "agent-ip": "192.168.1.2",
    "agent-port": "8080",
    "agent-sw-version": "3.1.0.1",
    "last-saved-configuration": "2017-10-18 - 07:24:16",
    "agent-up-time": 5350,
    "unsaved-config": 1
  },
  "id": 10
}
```

Chapter 5: Buffer Statistics and Tracking

5.1 Overview

The Buffer Statistics and Tracking (BST) feature provides a means to retrieve how each of the on-chip buffers are being utilized under different traffic scenarios. A collection of all such buffer-counts for various on-chip buffers is the BST report. This BST report is sent from the Agent to the collector either on-demand or periodically.

BST tracks either the current values or peak values of the buffer utilization. Also, BST can selectively track certain buffer categories.

BST allows thresholds to be set for some of the buffers and can send a notification to the client when the buffer usage exceeds the configured threshold.

BST allows clients to monitor for top ports suffering congestion or monitor the number of packets dropped due to congestion, on a per-port, per-queue-type (broadcast or multicast), or per-port-per-queue basis.

5.2 Buffer Utilization Statistics and Thresholds

Broadcom switch ASICs track buffer utilization for various on-chip buffers. The buffer utilization statistics are presented under the corresponding categories, called realms. To access (statistics of a) a buffer, the realm parameters must be provided. These parameters act as indices for the given buffer. For example, to access a buffer in the `egress-uc-queue` realm, the parameter is the corresponding queue ID. Similarly, to access a buffer in the `ingress-port-priority-group` realm, the parameters are the port ID and the priority-group ID. The device realm has no indices. No realm has more than two indices.

Each buffer, given its parameters, may have one or more statistic. For example, the `ingress-port-service-pool` realm buffer, for a given port and service pool ID, offers a single statistic: `um-share-buffer-count`. The `egress-port-service-pool` buffer, for a given port and service pool ID, offers four statistics: `uc-share-buffer-count`, `um-share-buffer-count`, `mc-share-buffer-count`, and `mc-share-queue-entries`. Each buffer statistic has an associated threshold for configuration. For example, for the `uc-share-buffer-count`, there is an associated `uc-share-threshold`.

Table 11 provides the list of realms and the associated indices for each realm. It also lists all available statistics and thresholds for each realm.

Table 11: Buffer Utilization Statistics and Thresholds^a

Realm	Index #1	Index #2	Statistic	Threshold
device	–	–	data	threshold
ingress-port-priority-group	port	priority-group	um-share-buffer-count, um-headroom-buffer-count	um-share-threshold, um-headroom-threshold
ingress-port-service-pool	port	service-pool	um-share-buffer-count	um-share-threshold
ingress-service-pool	service-pool	–	um-share-buffer-count	um-share-threshold
egress-port-service-pool	port	service-pool	uc-share-buffer-count, um-share-buffer-count, mc-share-buffer-count, mc-share-queue-entries	uc-share-threshold, um-share-threshold, mc-share-threshold, mc-share-queue-entries-threshold
egress-service-pool	service-pool	–	um-share-buffer-count, mc-share-buffer-count, mc-share-queue-entries	um-share-threshold, mc-share-threshold, mc-share-queue-entries-threshold
egress-uc-queue	queue	–	uc-buffer-count	uc-threshold
egress-uc-queue-group	queue-group	–	uc-buffer-count	uc-threshold
egress-mc-queue	queue	–	mc-buffer-count, mc-queue-entries	mc-threshold, mc-queue-entries-threshold
egress-cpu-queue	queue	–	cpu-buffer-count	cpu-threshold
egress-rqe-queue	queue	–	rqe-buffer-count	rqe-threshold

a. Note that not all counters/thresholds are supported on all silicon.

5.3 Configurability

The following configurability is provided by the BST feature:

- The BST feature can be enabled or disabled on the Agent.
- The BST feature can be configured to auto-accumulate the buffer statistics periodically and send the BST report to the collector at configurable intervals (as notification messages).
- The BST feature can provide buffer statistics (BST report) on demand.
- Buffer statistics collection (in the switch and in the Agent) of various subsets of buffer statistics (such as device-level, various ingress-groups, etc.) can be selectively disabled or enabled.
- The ASIC can be set up to gather the statistics of current values or peak values for the statistics.
- The Agent can be configured to report the buffer statistics in units of bytes, or in the units of cells, or in terms of a percentage of buffer usage (in terms of allocated levels).
- Various thresholds can be set up for the statistics in the ASIC, and the current values can be retrieved.
- The Agent can be set up to asynchronously send the buffer statistics report, called the *trigger report*, when the buffer usage exceeds the configured threshold.
- The BST feature can provide the number of packets dropped within the ASIC because of the congestion, for top ports suffering congestion, at a per-port, per-queue-type (broadcast or multicast), or per-port-per-queue level.

Any changes made to the feature configuration come into effect starting with the next reporting period.

5.4 Complete and Incremental Reports

Keeping in mind the extensive data that is likely to be collected and communicated, the following constraints are imposed on the BST reports.

- All of the on-demand BST reports are complete in nature.
- All of the asynchronous BST reports are incremental in nature, unless configured otherwise.
 - The BST feature compares the statistics collected in the current cycle to that of the previous cycle and marks the changed statistics in its native representation.
 - The JSON packing from the native representation is aware of the marking and ignores those statistics that have not changed since the last cycle.
 - Any statistics that are not supported by the underlying ASIC or the SDK are explicitly encoded with special value in the JSON message. These statistics are not included in subsequent asynchronous reports (being incremental).

5.5 Queues

The ASIC allows various on-chip queues to be assigned to individual ports. For example, a queue (say numbered 453) can be assigned as queue 4 to the port 45. Some of the ASICs allow any number of queues to be assigned to a port, whereas some have a fixed assignment. This assignment is user driven. Thus, the queues are identified in two forms: the asic-specified-queue-number, and the combination of user-assigned-queue-number and the port number.

The asic-specified-queue-number is referred to as hw-queue, or simply queue in this document and in the JSON messaging. The user assigned-queue-number is referred to as user-queue.

The APIs provided by the BST feature, which have the queue information specified in the response or request, allow both hw-queue and user-queue formats. More specifically:

- A majority of API requests that need the queue to be specified can accept *either of the formats*. The `configure-bst-multi-thresholds` API accepts the queue information user-queue format (i.e., along with the port number).
- The response/report messages contain the queue information in *both the formats*.

5.6 Deciphering Reports

While reporting statistics for a given realm, the statistics names are not included in the report. A JSON positional parameter method is used. In addition, one (and not more than one) of the indices is included as the first element in the array. This scheme is used to reduce the JSON message size for report messages. This scheme is applicable only for BST reports and threshold reports. For configuration messages to the Agent, all parameters are named JSON parameters.

This method is illustrated below with a few examples.

5.6.1 Example 1

Consider the following report snippet:

```
{
    "realm": "ingress-service-pool",
    "data": [[1, 3240], [2, 3660]]
}
```

The `ingress-service-pool` has a single index (`service pool`), and it is mentioned in the JSON array. See [Table 11, Buffer Utilization Statistics and Thresholds](#).

This report indicates the following.

- The realm is the `ingress-service-pool`.
- The service pool with ID 1 has the `uc-share-buffer-count` value of 3240.
- The service pool with ID 2 has the `uc-share-buffer-count` value of 3660.

5.6.2 Example 2

Consider the following report snippet:

```
{
  "realm": "ingress-port-priority-group",
  "data": [{
    "port": "2",
    "data": [[5, 45500, 44450]]
  }, {
    "port": "3",
    "data": [[5, 6700, 250], [7, 12667, 13456]]
  }]
}
```

The ingress-port-priority-group has two indices: port and the priority-group. See [Table 11, Buffer Utilization Statistics and Thresholds](#). The top-level index (port) is mentioned explicitly, and the second index (priority-group) is included as the first element in the JSON array.

This report indicates the following:

- The realm is the ingress-port-priority-group.
- The port 2 report has the following entries:
 - priority-group id 5 has
 - The um-share-buffer-count value is 45500.
 - The um-headroom-buffer-count value is 44450.
- The port 3 report has the following entries:
 - priority-group id 5 has
 - The um-share-buffer-count value is 6700.
 - The um-headroom-buffer-count value is 250.
 - priority-group id 7 has
 - The um-share-buffer-count value is 12267.
 - The um-headroom-buffer-count value is 13456.

5.6.3 Example 3

Consider the following report snippet:

```
{
    "realm": "egress-mc-queue",
    "data": [[1, "1", 1, 34, 89], [10, "2", 1, 1244, 0]]
}
```

The `egress-mc-queue` has a single index (queue). It is included as the first element in the JSON array. The second element is *port*, to which queue is assigned. The third element is the *user-queue*.

This report indicates the following:

- The realm is the `egress-mc-queue`.
- The Queue #1 (port #1 and user queue #1) has the following entries:
 - The `mc-buffer-count` value is 34.
 - The `mc-queue-entries` value is 89.
- The Queue #10 (port #2 and user queue #1) has the following entries:
 - The `mc-buffer-count` value is 1244.
 - The `mc-queue-entries` value is 0.

5.6.4 Example 4

Consider the following report snippet:

```
{
    "realm": "egress-uc-queue",
    "data": [[1, "1", 1, 34]]
}
```

The `egress-uc-queue` has a single index (queue). It is included as the first element in the JSON array. The second element is *port* to which queue is assigned. The third element is the *user-queue*.

This report indicates the following:

- The realm is the `egress-uc-queue`.
- The Queue #1 (port #1 and user queue #1) has the following entry
 - The `uc-buffer-count` value is 34.

The *port* is included as a second parameter for the buffer usage report and threshold reports for the `egress-uc-queue` and `egress-mc-queue` realms. Similarly the *user-queue* is included as the third parameter.

5.7 API

5.7.1 Overview

This section lists APIs supported by the BST feature. A description of each API is in subsequent sections.

Table 12: Buffer Statistics and Tracking APIs

API	Type	Description
configure-bst-feature	Configuration/Clear	Configure global options for BST and reporting.
configure-bst-tracking	Configuration/Clear	Setup tracking parameters.
configure-bst-thresholds	Configuration/Clear	Configure threshold/watermarks for a given buffer.
configure-bst-thresholds	Configuration/Clear	Configure multiple threshold/watermarks for specific buffers
clear-bst-statistics	Configuration/Clear	Clear all statistics.
clear-bst-thresholds	Configuration/Clear	Clear all configured thresholds.
get-bst-feature	Status/Reporting	Retrieve current BST configuration.
get-bst-tracking	Status/Reporting	Retrieve current tracking configuration.
get-bst-thresholds	Status/Reporting	Retrieve current threshold configuration.
get-bst-report	Status/Reporting	Obtain a snapshot/incremental buffer usage report.
Trigger Reports	Notification	Asynchronous report indicating a breach of a configured threshold.
get-bst-congestion-drop-counters	Status/Reporting	Obtain number of packets dropped due to congestion on a per-port/per-queue basis, immediately or periodically.
clear-bst-congestion-drop-counters	Configuration/Clear	Clear the congestion drop counters.
cancel-request^a	Configuration/Clear	Stop a previously initiated periodic reporting command.

a. This API is not specific to the BST feature and may be used to cancel any previously initiated periodic reporting command.

5.7.2 Configuration/Clear API

5.7.2.1 configure-bst-feature

The `configure-bst-feature` API configures the BST functionality on the Agent⁵. The parameters associated with this command are described in [Table 13](#).

Table 13: configure-bst-feature Parameters

Parameter	Type	Description
<code>bst-enable</code>	Boolean	Determines whether the BST feature should be active or not, on the Agent. By default, the BST feature is inactive on the Agent.
<code>send-async-reports</code>	Boolean	When enabled, the BST feature asynchronously collects the buffer statistics and sends the BST reports to the collector. By default, the asynchronous reporting is turned off.
<code>collection-interval</code>	Integer (range 0 to 600)	Determines the periodicity with which the BST reports are sent to the collector. The Agent may also use this value to read the statistics from the ASIC. The units for this parameter are seconds. The default value is 60 seconds.
<code>stats-in-percentage</code>	Boolean	When enabled, the Agent reports the buffer usage statistics as a percentage. When <code>stats-in-percentage</code> is enabled, the parameter <code>stat-units-in-cells</code> ^a is ignored while reporting the statistics. This field is applicable for statistics and threshold reporting. This is an optional parameter. By default, the value is set to false. The percentage value in the BST/trigger report is an approximation of buffer utilization, not an exact value.
<code>stat-units-in-cells</code>	Boolean	Determines whether the Agent reports the buffer statistics in the units of bytes or in the units of cells. Set to <i>false</i> , by default. The Agent reports the usage in bytes.
<code>trigger-rate-limit</code>	Integer (range 1 to 30)	Determines the maximum number of trigger reports for the configured interval (as specified by the <code>trigger-rate-limit-interval</code> parameter below) to be sent to the collector. This is an optional parameter, and the default value is 1.
<code>trigger-rate-limit-interval</code>	Integer (range 1 to 60)	Determines the interval during which the number of trigger-reports is rate limited (<code>trigger-rate-limit</code>) for sending to the collector. This is an optional parameter, and the default value is 1 second.
<code>send-snapshot-on-trigger</code>	Boolean	Determines whether the Agent should send a complete buffer statistics report for all configured realms to the collector, when a threshold is breached. If set to <i>true</i> , trigger-report contains buffer statistics for all configured realms. If set to <i>false</i> , trigger-report contains buffer statistics only for the statistic/ counter for which the trigger was raised. This is an optional parameter, and the default value is <i>true</i> .
<code>async-full-reports</code>	Boolean	When enabled, the BST feature asynchronously sends full BST reports to the collector. This is an optional parameter. By default, the asynchronous full reporting is turned off.

a. Cells are buffers inside the ASIC that are used to hold parts of a packet. Refer to the hardware documentation for details.

5. But not on the ASIC. See the next command.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "configure-bst-feature",
  "asic-id": "1",
  "params": {
    "bst-enable": 1,
    "send-async-reports": 1,
    "collection-interval": 300,
    "stat-units-in-cells": 0,
    "trigger-rate-limit": 5,
    "trigger-rate-limit-interval": 2,
    "send-snapshot-on-trigger": 1,
    "async-full-reports": 1,
    "stats-in-percentage": 0
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

5.7.2.2 configure-bst-tracking

The configure-bst-tracking command sets up the BST trackers and the tracking-mode on the ASIC. The parameters associated with the command are described in

Table 14: configure-bst-tracking Parameters

Parameter	Type	Description
track-peak-stats	Boolean	Determines whether ASIC tracks the current buffer usage count or the peak buffer usage. Set to <i>false</i> , by default. The ASIC tracks current buffer count.
track-ingress-port-priority-group	Boolean	When enabled, the ASIC actively tracks ingress per-port and per-priority group buffers. By default, this tracking is turned on.
track-ingress-port-service-pool	Boolean	When enabled, the ASIC actively tracks ingress per-port and per-service pool buffers. By default, this tracking is turned on.
track-ingress-service-pool	Boolean	When enabled, the ASIC actively tracks ingress per-service pool buffers. By default, this tracking is turned on.
track-egress-port-service-pool	Boolean	When enabled, the ASIC actively tracks egress per-port and per-service pool buffers. By default, this tracking is turned on.
track-egress-service-pool	Boolean	When enabled, the ASIC actively tracks egress per-service pool buffers. By default, this tracking is turned on.
track-egress-uc-queue	Boolean	When enabled, the ASIC actively tracks egress per unicast queue buffers. By default, this tracking is turned on.
track-egress-uc-queue-group	Boolean	When enabled, the ASIC actively tracks egress per unicast queue group buffers. By default, this tracking is turned on.
track-egress-mc-queue	Boolean	When enabled, the ASIC actively tracks egress per multicast queue buffers. By default, this tracking is turned on.

Table 14: configure-bst-tracking Parameters (Continued)

Parameter	Type	Description
track-egress-cpu-queue	Boolean	When enabled, the ASIC actively tracks egress per CPU queue buffers. By default, this tracking is turned on.
track-egress-rqe-queue	Boolean	When enabled, the ASIC actively tracks egress per RQE queue buffers. By default, this tracking is turned on.
track-device	Boolean	When enabled, the ASIC actively tracks per-device (ASIC) buffers. By default, this tracking is turned on.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "configure-bst-tracking",
  "asic-id": "1",
  "params": {
    "track-peak-stats" : 1,
    "track-ingress-port-priority-group" : 1,
    "track-ingress-port-service-pool" : 1,
    "track-ingress-service-pool" : 1,
    "track-egress-port-service-pool" : 1,
    "track-egress-service-pool" : 1,
    "track-egress-uc-queue" : 1,
    "track-egress-uc-queue-group" : 1,
    "track-egress-mc-queue" : 1,
    "track-egress-cpu-queue" : 1,
    "track-egress-rqe-queue" : 1,
    "track-device" : 1
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

5.7.2.3 configure-bst-thresholds

The `configure-bst-thresholds` command sets up the BST thresholds for various realms in the ASIC.

When a threshold is configured, and the associated buffer usage exceeds the threshold value, then an asynchronous notification is sent by the Agent.

The parameters associated with this command depend on the realm for which the thresholds are being set up. The list of realms and associated parameters/thresholds for each realm is provided in [Section 5.2, Buffer Utilization Statistics and Thresholds](#).

At any point in time, the valid set of thresholds depends on the MMU configuration of the Broadcom ASIC. Consult the Broadcom ASIC documentation for supported combinations.

Threshold configurations are in terms of number of bytes, number of cells, or percentage of allocated buffer size. The actual unit is determined by the Agent based on its current configuration for reporting the statistics set via `configure-bst-feature` command. For example, when the Agent is configured to report statistics in terms of percentage via `configure-bst-feature` command, then the Agent accepts threshold configuration in terms of percentage.

Sample JSON-RPC requests are shown below.

```
{
  "jsonrpc": "2.0",
  "method": "configure-bst-thresholds",
  "asic-id": "1",
  "params": {
    "realm": "egress-rqe-queue",
    "queue": 1,
    "rqe-threshold": 15156
  },
  "id": 1
}

{
  "jsonrpc": "2.0",
  "method": "configure-bst-thresholds",
  "asic-id": "1",
  "params": {
    "realm": "egress-port-service-pool",
    "port": "1",
    "service-pool": 3,
    "uc-share-threshold": 15156,
    "um-share-threshold": 15156,
    "mc-share-threshold": 15156,
    "mc-share-queue-entries-threshold": 15156
  },
  "id": 1
}

{
  "jsonrpc": "2.0",
  "method": "configure-bst-thresholds",
  "asic-id": "1",
  "params": {
    "realm": "egress-mc-queue",
    "port": "105",
    "user-queue": 0,
```

```

    "mc-threshold": 100
  },
  "id": 5
}

```

The Agent acknowledges the command and returns an appropriate error code if needed.

5.7.2.4 configure-bst-multi-thresholds

The `configure-bst-multi-thresholds` command sets up one or more BST thresholds for various realms in the ASIC. The threshold configuration behavior is identical to the `configure-bst-thresholds` command.

The parameters associated with this command depend on the realm for which the thresholds are being set up. The list of realms and associated parameters/thresholds for each realm is provided in [Section 5.2, Buffer Utilization Statistics and Thresholds](#).

At any point in time, the valid set of thresholds depends on the MMU configuration of the Broadcom ASIC. Consult the Broadcom ASIC documentation for supported combinations.

Threshold configurations are in terms of number of bytes, number of cells, or percentage of allocated buffer size. The actual unit is determined by the Agent based on its current configuration for reporting the statistics set via the `configure-bst-feature` command. For example, when the Agent is configured to report statistics in terms of percentage via the `configure-bst-feature` command, then the Agent accepts threshold configuration in terms of percentage.

Sample JSON-RPC requests are shown below.

```

{
  "jsonrpc": "2.0",
  "method": "configure-bst-multi-thresholds",
  "asic-id": "1",
  "params": [{
    "realm": "device",
    "data": 46
  }, {
    "realm": "ingress-port-priority-group",
    "data": [{
      "port": "2",
      "data": [[5, 45500, 44450]]
    }, {
      "port": "3",
      "data": [[5, 45500, 44450]]
    }]
  }, {
    "realm": "ingress-port-service-pool",
    "data": [{
      "port": "2",
      "data": [[5, 324]]
    }, {
      "port": "3",
      "data": [[6, 366]]
    }]
  }, {
    "realm": "ingress-service-pool",
    "data": [[1, 3240], [2, 3660]]
  }, {
    "realm": "egress-cpu-queue",
    "data": [[3, 4566]]
  }
]
}

```



```

    }, {
      "realm": "egress-mc-queue",
      "data": [[ "1", 1, 34, 89], [ "4", 2, 1244, 0], [ "5", 1, 0, 3]]
    }, {
      "realm": "egress-port-service-pool",
      "data": [{
        "port": "2",
        "data": [[5, 0, 324, 0]]
      }, {
        "port": "3",
        "data": [[6, 0, 366, 0]]
      }]
    }, {
      "realm": "egress-rqe-queue",
      "data": [[2, 3333], [5, 25]]
    }, {
      "realm": "egress-service-pool",
      "data": [[2, 0, 0, 3240], [3, 3660, 0, 0]]
    }, {
      "realm": "egress-uc-queue",
      "data": [[ "1", 0, 1111]]
    }, {
      "realm": "egress-uc-queue-group",
      "data": [[6, 2222]]
    }],
    "id": 1
  }
}

```

In addition to configuring values for the thresholds, this API can also be used to clear the existing threshold by passing 0 as the value for the threshold. This resets the configured threshold to the silicon default value (no threshold).

A value of -1 for the threshold for a counter results in no change to the current setting for the associated threshold in the silicon. This is useful in cases where the realms have multiple counters and the user may prefer to set the threshold for only one counter but not change others.

The Agent acknowledges the command and returns an appropriate error code if needed.

5.7.2.5 clear-bst-statistics

The `clear-bst-statistics` command clears all the BST statistics on the Agent⁶. This API does not require any parameters.

A sample JSON-RPC request is shown below.

```

{
  "jsonrpc": "2.0",
  "method": "clear-bst-statistics",
  "asic-id": "1",
  "params": {
  },
  "id": 1
}

```

The Agent acknowledges the command and returns an appropriate error code if needed.

6. A side effect of this command is that the following asynchronous BST report will be complete, including all realms.

5.7.2.6 clear-bst-thresholds

The `clear-bst-thresholds` command clears all configured BST thresholds on the Agent. This API does not require any parameters.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "clear-bst-thresholds",
  "asic-id": "1",
  "params": {
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

Instead of clearing all thresholds, clearing a specific threshold can be achieved by setting the corresponding threshold to 0 by using the `configure-bst-multi-thresholds` command.

5.7.2.7 clear-bst-congestion-drop-counters

The `clear-bst-congestion-drop-counters` command clears all the Congestion Drop counters on the Silicon. This API does not require any parameters.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "clear-bst-congestion-drop-counters",
  "asic-id": "1",
  "params": {
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

5.7.3 Status/Reporting API

5.7.3.1 get-bst-feature

The get-bst-feature command is used to retrieve the current configuration of the BST functionality on the Agent. This API does not require any parameters.

For the response, the Agent returns the same parameters as sent via the configure-bst-feature command.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-feature",
  "asic-id": "1",
  "params": { },
  "id": 1
}
```

An associated sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-feature",
  "asic-id": "1",
  "version": "1",
  "result": {
    "bst-enable": 1,
    "send-async-reports": 0,
    "collection-interval": 200,
    "stats-in-percentage": 0,
    "stat-units-in-cells": 0,
    "trigger-rate-limit": 0,
    "trigger-rate-limit-interval": 0,
    "send-snapshot-on-trigger": 0
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

5.7.3.2 get-bst-tracking

The `get-bst-tracking` command is used to retrieve the current BST tracking configuration of the Broadcom ASIC. This API does not require any parameters.

For the response, the Agent returns the same parameters as those sent via the `configure-bst-tracking` command.

This command is valid even when the feature is disabled on the Agent.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-tracking",
  "asic-id": "1",
  "params": { },
  "id": 1
}
```

An associated sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-tracking",
  "asic-id": "1",
  "version": "1",
  "result": {
    "track-peak-stats" : 1,
    "track-ingress-port-priority-group" : 1,
    "track-ingress-port-service-pool" : 1,
    "track-ingress-service-pool" : 1,
    "track-egress-port-service-pool" : 1,
    "track-egress-service-pool" : 1,
    "track-egress-uc-queue" : 1,
    "track-egress-uc-queue-group" : 1,
    "track-egress-mc-queue" : 1,
    "track-egress-cpu-queue" : 1,
    "track-egress-rqe-queue" : 1,
    "track-device" : 1      },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

5.7.3.3 get-bst-thresholds

The `get-bst-thresholds` command is used to retrieve the currently configured BST thresholds from the Agent. The collector may choose to request a selective report (see parameters below).

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-thresholds",
  "asic-id": "1",
  "params": {
    "include-ingress-port-priority-group" : 1,
    "include-ingress-port-service-pool" : 1,
    "include-ingress-service-pool" : 1,
    "include-egress-port-service-pool" : 1,
    "include-egress-service-pool" : 1,
    "include-egress-uc-queue" : 1,
    "include-egress-uc-queue-group" : 1,
    "include-egress-mc-queue" : 1,
    "include-egress-cpu-queue" : 1,
    "include-egress-rqe-queue" : 1,
    "include-device" : 1
  },
  "id": 1
}
```

The Agent always returns a complete report for the requested thresholds.

The thresholds in the ASIC are set by the Agent only as part of the `configure-bst-thresholds` command. Unlike statistics, the threshold configuration does not change on its own. For this reason, it is not necessary to retrieve a periodic threshold report from the Agent.

Threshold reporting is in terms of bytes, cells, or percentage of maximum allocated buffer size. The actual reporting unit is determined by the Agent based on its current configuration for reporting the statistics via the `configure-bst-feature` command. For example, when the Agent is configured to report statistics in terms of percentage via the `configure-bst-feature` command, the Agent reports thresholds in terms of percentage.

The `method` object of the JSON message in the BST report is set to `get-bst-thresholds` for the response message.

A sample BST threshold report indicating various realms (categories) and the associated data is shown below. It is provided for illustrative purposes only and does not include all port/queue/pool data in it. Rather, the JSON array is used to indicate the multiplicity while providing data for a single element.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-thresholds",
  "asic-id": "20",
  "version": "1",
  "time-stamp": "2014-11-14 - 00:15:04 ",
  "report": [{
    "realm": "device",
    "data": 46
  }, {
    "realm": "ingress-port-priority-group",
    "data": [{
      "port": "2",
      "data": [[5, 45500, 44450]]
    }]
  }]
}
```

```
    }, {
      "port": "3",
      "data": [[5, 45500, 44450]]
    }],
  }, {
    "realm": "ingress-port-service-pool",
    "data": [{
      "port": "2",
      "data": [[5, 324]]
    }, {
      "port": "3",
      "data": [[6, 366]]
    }
  ]},
  {
    "realm": "ingress-service-pool",
    "data": [[1, 3240], [2, 3660]]
  }, {
    "realm": "egress-cpu-queue",
    "data": [[3, 4566, 0]]
  }, {
    "realm": "egress-mc-queue",
    "data": [[1, "1", 1, 34, 89], [2, "4", 2, 1244, 0], [3, "5", 1, 0, 3]]
  }, {
    "realm": "egress-port-service-pool",
    "data": [{
      "port": "2",
      "data": [[5, 0, 324, 0]]
    }, {
      "port": "3",
      "data": [[6, 0, 366, 0]]
    }
  ]},
  {
    "realm": "egress-rqe-queue",
    "data": [[2, 3333, 4444], [5, 25, 45]]
  }, {
    "realm": "egress-service-pool",
    "data": [[2, 0, 0, 3240], [3, 3660, 0, 0]]
  }, {
    "realm": "egress-uc-queue",
    "data": [[6, "1", 0, 1111]]
  }, {
    "realm": "egress-uc-queue-group",
    "data": [[6, 2222]]
  }
]
```

5.7.3.4 get-bst-report

The `get-bst-report` command is used to retrieve the current BST report from the Agent. The collector may choose to request for selective report (see parameters in [Table 15](#)).

Table 15: get-bst-report Parameters

Parameter	Type	Description
<code>include-ingress-port-priority-group</code>	Boolean	When set, the Agent includes the ingress per-port per-priority group buffer statistics into the report.
<code>include-ingress-port-service-pool</code>	Boolean	When set, the Agent includes the ingress per-port per-service pool buffer statistics into the report.
<code>include-ingress-service-pool</code>	Boolean	When set, the Agent includes the ingress per-service pool buffer statistics into the report.
<code>include-egress-port-service-pool</code>	Boolean	When set, the Agent includes the egress per-port per-service pool buffer statistics into the report.
<code>include-egress-service-pool</code>	Boolean	When set, the Agent includes the egress per-service pool buffer statistics into the report.
<code>include-egress-uc-queue</code>	Boolean	When set, the Agent includes the egress per-unicast queue buffer statistics into the report.
<code>include-egress-uc-queue-group</code>	Boolean	When set, the Agent includes the egress per-unicast queue group buffer statistics into the report.
<code>include-egress-mc-queue</code>	Boolean	When set, the Agent includes the egress per-multicast queue buffer statistics into the report.
<code>include-egress-cpu-queue</code>	Boolean	When set, the Agent includes the egress per-CPU queue buffer statistics into the report.
<code>include-egress-rqe-queue</code>	Boolean	When set, the Agent includes the ingress per-RQE queue buffer statistics into the report.
<code>include-device</code>	Boolean	When set, the Agent includes the ingress device level buffer statistics into the report.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-report",
  "asic-id": "1",
  "params": {
    "include-ingress-port-priority-group" : 1,
    "include-ingress-port-service-pool" : 1,
    "include-ingress-service-pool" : 1,
    "include-egress-port-service-pool" : 1,
    "include-egress-service-pool" : 1,
    "include-egress-uc-queue" : 1,
    "include-egress-uc-queue-group" : 1,
    "include-egress-mc-queue" : 1,
    "include-egress-cpu-queue" : 1,
    "include-egress-rqe-queue" : 1,
    "include-device" : 1
  },
  "id": 1
}
```

The Agent always returns a complete report for the requested buffer statistics.

It must be noted that the Agent returns the BST report following the command. However, the asynchronous reporting cycle may be reset.

The `method` object of the JSON message in the BST report is set to `get-bst-report` for both the asynchronous notifications as well as the response message.

The parameters associated with this command depend on the realm for which the buffer usage (count) is being reported. The list of realms and associated parameters / buffer-count for each realm is provided in [Section 5.2, Buffer Utilization Statistics and Thresholds](#).

It may be noted that at any point in time, the valid set of buffer counts (usage) depends on the MMU configuration of the Broadcom ASIC. Refer to the Broadcom ASIC documentation for supported configurations.

A sample BST report indicating various realms (categories) and the associated statistics is shown below. It is provided for illustrative purposes only and does not include all the realm/port/queue/pool data in it. Rather, the JSON array is used to indicate the multiplicity while providing data for a single element.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-report",
  "asic-id": "20",
  "version": "1",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "report": [{
    "realm": "device",
    "data": 46
  }, {
    "realm": "ingress-port-priority-group",
    "data": [{
      "port": "2",
      "data": [[5, 45500, 44450]]
    }, {
      "port": "3",
      "data": [[5, 45500, 44450]]
    }
  ]
}, {
  "realm": "ingress-port-service-pool",
  "data": [{
    "port": "2",
    "data": [[5, 324]]
  }, {
    "port": "3",
    "data": [[6, 366]]
  }
]
}, {
  "realm": "ingress-service-pool",
  "data": [[1, 3240], [2, 3660]]
}, {
  "realm": "egress-cpu-queue",
  "data": [[3, 4566, 0]]
}, {
  "realm": "egress-mc-queue",
  "data": [[1, "1", 1, 34, 89], [2, "4", 2, 1244, 0], [3, "5", 2, 0, 3]]
}, {
  "realm": "egress-port-service-pool",
  "data": [{
    "port": "2",
```



```

        "data": [[5, 0, 324, 0]]
      }, {
        "port": "3",
        "data": [[6, 0, 366, 0]]
      }
    ], {
      "realm": "egress-rqe-queue",
      "data": [[2, 3333, 4444], [5, 25, 45]]
    }, {
      "realm": "egress-service-pool",
      "data": [[2, 0, 0, 3240], [3, 3660, 0, 0]]
    }, {
      "realm": "egress-uc-queue",
      "data": [[6, "1", 2, 1111]]
    }, {
      "realm": "egress-uc-queue-group",
      "data": [[6, 2222]]
    }
  ]
}

```

5.7.3.5 get-bst-congestion-drop-counters

The `get-bst-congestion-drop-counters` command is used to retrieve the congestion drop counters periodically or immediately from the Agent. There are different sets of drop counters that can be retrieved by the collector. Based on the type of the set, the number of parameters passed to the command varies. The basic parameters are provided in [Table 16](#).

Table 16: get-bst-congestion-drop-counters Parameters

Parameter	Type	Description
request-type	String	Indicates the specific set of drop counters being requested. The following sets are supported: <ul style="list-style-type: none"> ■ top-drops. Ports suffering maximum congestion in the switch and the associated drop counters. ■ top-port-queue-drops. Top port-queue level drop-counters in the switch. ■ port-drops. Per-port total drop counters. ■ port-queue-drops. Port-queue level drop-counters.
collection-interval	Integer (range 0 to 3600)	Determines the period with which the congestion drop counters are collected from the ASIC and reported to the client. The units for this parameter are seconds. This is an optional parameter. The default value is 0 (zero), which indicates an immediate response.

[Table 17](#) shows drop report parameters for the top congestion ports/port-queues report.

Table 17: Top Congestion Ports/Port-Queues Drop Report Parameters

Parameter	Type	Description
count	Integer (range 0 to 64)	Number of ports required in the report. The ports are placed in a sorted list, with the port suffering maximum congestion at the top. Out of this list, the count number of ports and their drop-counters are reported back to the client.
queue-type	String	This represents the queue type filter for the report. Possible values are: <ul style="list-style-type: none"> ■ ucast. Indicates unicast queues. ■ mcast. Indicates multicast queues. ■ all. Indicates all supported queue types in the ASIC.

Drop report parameters for (all/specific) congestion ports/port-queues report are listed in [Table 18](#).

Table 18: All/Specific Congestion Ports/Port-Queues Drop Report Parameters

Parameter	Type	Description
port-list	Array	Comma separated list of ports for which congestion drop counter report is requested. A special keyword <i>all</i> is allowed to get congestion drop counters for all ports.
queue-type	String	This represents the queue type filter for the report. Possible values are: <ul style="list-style-type: none">■ ucast. Indicates unicast queues.■ mcast. Indicates multicast queues.■ all. Indicates all supported queue types in the ASIC.
queue-list	Array	An array of queue numbers to be considered for the drop report.

A sample JSON-RPC request to get the top eight congestion drop counters per port is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-congestion-drop-counters",
  "asic-id": "1",
  "params": {
    "request-type" : "top-drops",
    "request-params": {
      "count": 8
    },
    "collection-interval": 30
  },
  "id": 1
}
```

A sample BST congestion drop report for the above request is as follows:

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-congestion-drop-counters",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2016-1-1 - 00:15:04 ",
  "report": [{
    "report-type": "top-drops",
    "data": [{
      "port": "2",
      "data": 8500
    }, {
      "port": "8",
      "data": 7550
    }, {
      "port": "6",
      "data": 6500
    }, {
      "port": "1",
      "data": 5550
    }, {
      "port": "7",
      "data": 4500
    }, {
      "port": "10",
      "data": 3550
    }, {
      "port": "20",
      "data": 2500
    }, {
      "port": "19",
      "data": 1550
    }
  ]
}, {
  "id": 1
}]
}
```

A sample JSON-RPC request to get the top eight congestion drop counters per port-queue is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-congestion-drop-counters",
  "asic-id": "1",
  "params": {
    "request-type" : "top-port-queue-drops",
    "request-params": {
      "count": 8,
      "queue-type": "all"
    },
    "collection-interval": 30
  },
  "id": 1
}
```

A sample BST congestion drop report for the above request is as follows:

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-congestion-drop-counters",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2016-1-1 - 00:15:04 ",
  "report": [{
    "report-type": "top-port-queue-drops",
    "data": [{
      "port": "2",
      "queue-type": "ucast",
      "data": [[1,8500],[4,8400]]
    }, {
      "port": "8",
      "queue-type": "mcast",
      "data": [[4,7550]]
    }, {
      "port": "6",
      "queue-type": "ucast",
      "data": [[2,6500],[1,5400]]
    }, {
      "port": "2",
      "queue-type": "mcast",
      "data": [[3,4500],[1,4400],[5,200]]
    }
  ]
}, {
  "id": 1
}
```

A sample JSON-RPC request to get all/specific port congestion drop counter information is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-congestion-drop-counters",
  "asic-id": "1",
  "params": {
    "request-type": "port-drops",
    "request-params": {
      "port-list": ["1", "2", "3", "4"]
    },
    "collection-interval": 30
  },
  "id": 1
}
```

A sample BST congestion drop report for the above request is as follows:

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-congestion-drop-counters",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2016-1-1 - 00:15:04 ",
  "report": [{
    "report-type": "port-drops",
    "data": [{
      "port": "1",
      "data": 8500
    }, {
      "port": "2",
      "data": 27550
    }, {
      "port": "3",
      "data": 16500
    }, {
      "port": "4",
      "data": 5550
    }
  ]
}],
  "id": 1
}
```

A sample JSON-RPC request to get all port-queue/ specific port-queue congestion drop counter information is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-congestion-drop-counters",
  "asic-id": "1",
  "params": {
    "request-type": "port-queue-drops",
    "request-params": {
      "port-list": ["1", "2"],
      "queue-type": "all",
      "queue-list": [1, 2, 3]
    },
    "collection-interval": 30
  },
  "id": 1
}
```

A sample BST congestion drop report for the above request is as follows:

```
{
  "jsonrpc": "2.0",
  "method": "get-bst-congestion-drop-counters",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2016-1-1 - 00:15:04 ",
  "report": [{
    "report-type": "port-queue-drops",
    "data": [{
      "port": "1",
      "queue-type": "ucast",
      "data": [[1,8500],[2,8400], [3,156000]]
    }, {
      "port": "1",
      "queue-type": "mcast",
      "data": [[1,600],[2,400], [3,1000]]
    }, {
      "port": "2",
      "queue-type": "ucast",
      "data": [[1,500],[2,800], [3,56000]]
    }, {
      "port": "2",
      "queue-type": "mcast",
      "data": [[1,2600],[2,3400], [3,21000]]
    }
  ]
}, {
  "id": 1
}
```

The Agent returns the requested drop counter report immediately. When the `collection-interval` is non-zero, the Agent continues to send requested drop counter reports at the configured interval asynchronously as notifications.

The Agent can be requested to serve multiple periodic reports.

The Agent acknowledges the command and returns an appropriate error code if needed.

5.7.4 Trigger Reports

When any of the configured threshold (via the `configure-bst-thresholds`) API is triggered, that is, the buffer usage crosses the configured threshold, an asynchronous notification is sent by the Agent. This notification is identical to an asynchronous BST report with the following two exceptions:

- The `method` object has the value of `trigger-report`.
- Additional fields are included to specify the actual counter causing the threshold breach.

A partial sample report is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "trigger-report",
  "asic-id": "20",
  "version": "1",
  "time-stamp": "2014-11-18 - 00:13:08 ",
  "realm": "ingress-port-priority-group",
  "counter": "um-share-buffer-count",
  "port": "2",
  "priority-group": "5",
  "report": [{
    "realm": "device",
    "data": 46
  }, {
    "realm": "ingress-port-priority-group",
    "data": [{
      "port": "2",
      "data": [[5, 45500, 44450]]
    }, {
      "port": "3",
      "data": [[5, 45500, 44450]]
    }]
  }, {
    "realm": "ingress-port-service-pool",
    "data": [{
      "port": "2",
      "data": [[5, 324]]
    }, {
      "port": "3",
      "data": [[6, 366]]
    }]
  }, {
    "realm": "ingress-service-pool",
    "data": [[1, 3240], [2, 3660]]
  }, {
    "realm": "egress-cpu-queue",
    "data": [[3, 4566, 0]]
  }, {
    "realm": "egress-mc-queue",
    "data": [[1, "1", 1, 34, 89], [2, "4", 2, 1244, 0], [3, "5", 1, 0, 3]]
  }, {
    "realm": "egress-port-service-pool",
    "data": [{
      "port": "2",
      "data": [[5, 0, 324, 0]]
    }, {
      "port": "3",
```

```

        "data": [[6, 0, 366, 0]]
      ]]
    }, {
      "realm": "egress-rqe-queue",
      "data": [[2, 3333, 4444], [5, 25, 45]]
    }, {
      "realm": "egress-service-pool",
      "data": [[2, 0, 0, 3240], [3, 3660, 0, 0]]
    }, {
      "realm": "egress-uc-queue",
      "data": [[6, "1", 1, 1111]]
    }, {
      "realm": "egress-uc-queue-group",
      "data": [[6, 2222]]
    }
  ]]
}

```

NOTE: When a trigger situation is encountered (that is, any of the configured thresholds are breached), the ASIC freezes all the BST counters until the feature is re-enabled/re-configured. The Agent re-enables the BST on the ASIC upon a trigger notification and sends the trigger report to the client. The number of triggers in a given duration can be rate-limited via the `configure-bst-feature` command.

Chapter 6: Packet Trace

6.1 Overview

When a packet enters a switch, it is processed and forwarded out on one or more destination ports. Typically, the results of intermediate packet processing steps are not visible. The Packet Trace feature provides detailed information on these intermediate processes, like egress hashing information for LAG and ECMP. This detailed information is called a *trace-profile*.

The Packet Trace feature allows the client to inject a packet into the ingress packet processing pipeline that is then processed as if it were received on one of the front panel ports. The Packet Trace feature then logs the internal forwarding states. This can be useful in the diagnosis of unexpected errors or as an offline debugging tool.

The Packet Trace feature also provides an option to the client to match incoming live traffic packets against a client-specified match criterion and provide a *trace-profile* for such packets. In this case, the client can specify the packet match criterion instead of providing a packet to be injected.

In both the scenarios, the packet processing information—more specifically the egress port information for the packet—is provided back to the client.

Additionally, the Packet Trace feature allows clients to keep track of dropped packets and understanding the reasons for certain packet drops in the ASIC.

6.2 Configurability

The following configurability is provided by the Packet Trace feature.

- The Packet Trace feature can be enabled or disabled on the Agent.
- The Packet Trace feature can be configured to collect trace-profile periodically for a client supplied packet and send trace-profile to the client at configurable interval (as notification reports).
- The Packet Trace feature can provide a trace-profile for a client-supplied packet on demand.
- The Packet Trace feature can be configured to drop or forward the packet injected by requester after the trace-profile is collected by the Agent.
- The Packet Trace feature can be configured to collect a trace-profile for all live traffic packets on the wire, matching user-supplied criteria and send a trace-profile to the client (as notification reports).
- The Packet Trace feature can provide drop reason for a client-supplied packet or for live traffic packets on the wire, matching user-supplied criteria and send a drop-profile to the client (as notification reports).
- The Packet Trace feature can provide a count of packets dropped by the Silicon for live traffic packets received from the wire, matching user-supplied criteria and send a drop-profile to the client (as notification reports).
- All packet trace/drop profile requests that have been initiated can be canceled via the cancel-request API.

6.3 Packet Format

The BroadView Agent expects the client-supported packet to be in a PCAP format, which is a binary format. However, JSON has no data type to represent the binary data. This means that the collector must encode the PCAP in base64 format. The Agent decodes the PCAP information, which is in a base64 format, as binary and strips the PCAP global header and PCAP packet header before injecting the packet.

Each PCAP file can have multiple packets, but the BroadView Agent supports only one packet in a given base64-encoded PCAP data.

After decoding, the Agent checks the first four bytes for the magic number⁷. If the first four bytes is not the magic number, the input file data is ignored. The Agent also checks for the major and minor numbers. The supported major number is 2, and the supported minor number is 4. If the magic number and major and minor numbers are valid, then the Agent extracts the packet by removing the PCAP file header and the PCAP packet header.

Following is an example of the decoded packet (after base64 decoding) on the Agent.

```
a1 b2 c3 d4 0 2 0 4 0 0 0 0 0 0 0
0 0 ff ff 0 0 0 1 55 18 f3 f4 0 b 94 ea
0 0 0 78 0 0 0 78 11 22 33 44 55 66 77 88
99 aa bb cc 81 0 0 1 8 0 45 2 0 66 7f af
d3 a0 40 6 40 59 1 2 3 4 5 6 7 8 10 1
d2 24 10 2 2a 4 50 0 16 d0 5f eb d3 d0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 c 8a 6c f0
```

The Agent deciphers this packet as follows:

- The first four bytes are the magic number (a1 b2 c3 d4).
- The major version is 2 (0 2), which is two bytes.
- The minor version number (0 4) is two bytes.
- The total length of the PCAP file header is 24 bytes, and the PCAP packet length is 16 bytes.
- The four bytes from offset 32 indicate the captured packet length (00 00 00 78), and the subsequent four bytes (00 00 00 78) indicate the length of the packet on the wire.

If the captured packet length is smaller than the packet length on the wire, then the input is discarded, and an appropriate error code is returned by the Agent.

When the Agent sends a packet to the client, as part of the JSON response, the packet is formatted as a base-64 encoded raw packet (as received from the wire). The packets in the JSON response are not PCAP encoded packets.

In this document, wherever a packet is used in a JSON request/response, a partial encoded packet with trailing dots, as shown below, is used for brevity.

```
"packet": "0010203232.."
```

7. For more information about the PCAP magic number, refer to PCAP documentation available on the Internet.

6.4 Live Traffic Triggered Trace-Profile/Drop-Profile

Instead of the client providing a packet and requesting a trace-profile, the client can provide 5-tuple information and request a trace-profile for packets matching the specified criteria. The Agent copies the packets matching the conditions specified in the 5-tuple to the CPU, and then injects them to the switch ASIC to retrieve the packet trace profile. If this feature is supported, then the string `Live-PT` is included in the response to the API `get-switch-properties`.

The same 5-tuple definition can be used by the clients to request for a drop-profile for matching live traffic.

The 5-tuple parameters are listed in [Table 19](#).

Table 19: 5-Tuple Parameters

Parameter	Type	Description
<code>src-ip</code>	String	IP address is used to match the source IP address of the packet.
<code>dst-ip</code>	String	IP address is used to match destination IP address of the packet.
<code>protocol</code>	String	Protocol number is used to match layer-4 protocol number in the packet.
<code>l4-src-port</code>	String	Layer-4 source port number is used to match the L4 source port of the packet.
<code>l4-dst-port</code>	String	Layer-4 destination port number is used to match the L4 destination port of the packet.

The client may specify a wild card for any of these parameters. A wild card indicates a “match any” condition for that parameter. A wild card is specified using the string *any* in the JSON request.

Up to three wildcards in a given 5-tuple are accepted by the Agent for a trace-profile. For a drop-profile, all parameters can be specified as wildcards.

Trace/Drop profile information is sent asynchronously by the Agent for traffic-triggered trace-profile/drop-profile requests.

6.5 Timestamps

The trace-profile data contains the parameters `packet-received-time-stamp`, `packet-ingress-time-stamp` and `packet-egress-time-stamp`.

- The `packet-received-time-stamp` indicates the CPU time when the packet is received.
- The `packet-ingress-time-stamp` indicates the ASIC time when the packet is received by the ASIC.
- The `packet-egress-time-stamp` indicates the ASIC time when the packet leaves the ASIC.

The ability for the Agent to report either/both of the timestamps, `packet-ingress-time-stamp` and `packet-egress-time-stamp`, is determined by the capability of the ASIC. If the ASIC does not provide the necessary data, the Agent does not include those fields in the response JSON.

Both of the timestamps, `packet-ingress-time-stamp` and `packet-egress-time-stamp`, are valid only for live traffic triggered trace-profile reports.

6.6 Deciphering Trace-Profile

The ASIC reports the results of the packet trace (trace-profile) for different switching and routing tables. The results are presented under various categories, called *realms*.

[Table 20](#) provides the list of realms and the associated results for each realm. It also lists the available fields/information for each realm.

Table 20: Deciphering Trace-Profile Realms

Realm	Index	Trace-Profile Parameters
lag-link-resolution	—	lag-id, lag-members, dst-lag-member, fabric-lag-id ^a , fabric-lag-members, fabric-lag-dst-member
ecmp-link-resolution	ecmp-level	ecmp--group-id, ecmp-members, ecmp-dst-member, ecmp--dst-port, ecmp--next-hop-ip
link-resolution	—	dst-port

- a. The three parameters `fabric-lag-id`, `fabric-lag-members`, and `fabric-lag-dst-member` are deprecated, and the trace-profile reports sent by the Agent do not hold any value in these parameters. These are kept in the JSON structure for maintaining backward compatibility and may be removed in a future version.

The LAG and ECMP resolution parameters are standard parameters and need no further explanation. The `dst-port` parameter in the link-resolution realm indicates the egress-port on which the packet exited the switch.

A sample trace-profile report is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-profile",
  "asic-id": "1",
  "version": "1",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "report": [
    {
      "port": "1",
      "trace-profile": [
        {
          "realm": "lag-link-resolution",
          "data": {
            "lag-id": "2",
            "lag-members": ["1", "2", "3", "4"],
            "dst-lag-member": "4"
          }
        },
        {
          "realm": "ecmp-link-resolution",
          "data": [
            {
              "ecmp-group-id": "200256",
              "ecmp-members": [
                ["100005", "3.3.3.2", "15"],
                ["100006", "4.4.4.2", "16"]
              ],
              "ecmp-dst-member": "100005",
              "ecmp-dst-port": "15",
              "ecmp-next-hop-ip": "3.3.3.2"
            }
          ]
        }
      ]
    },
    {
      "port": "2",
      "trace-profile": [
        {
          "realm": "lag-link-resolution",
          "data": {
            "lag-id": "2",
            "lag-members": ["1", "2", "3", "4"],
            "dst-lag-member": "4"
          }
        },
        {
          "realm": "ecmp-link-resolution",
          "data": [
            {
              "ecmp-group-id": "200256",
              "ecmp-members": [
                ["100005", "3.3.3.2", "15"],
                ["100006", "4.4.4.2", "16"]
              ],
              "ecmp-dst-member": "100005",
              "ecmp-dst-port": "15",
            }
          ]
        }
      ]
    }
  ]
}
```

```

    "ecmp-next-hop-ip": "3.3.3.2"
  }
]
}
]
}
],
"id": 1
}

```

6.7 Deciphering Drop-Profile

The Drop-Profile is a combination of the following fields describing why the packet may have been dropped in the ASIC.

Parameter	Type	Description
drop-reason	Array of Strings	A list of human-readable texts, indicating the specific reason the packet may have been dropped in the ASIC. Note that there may be more than one reason for dropping a packet in the ASIC.
drop-reason-code	String	A hexadecimal value encoded as a String which indicates the drop-code. ^a
drop-location	Enumeration	Indicates the specific place in the ASIC the packet may have been dropped. There are three possible values for this parameter: dropped-at-ingress dropped-in-pipeline dropped-at-egress
port	String	Indicates the ingress port for the packet
packet	base64-encoded packet	Indicates the packet for this drop-profile.

a. This value is intended for communicating with Support while debugging a specific packet drop issue.

A sample drop-profile report is shown below.

```

{
  "jsonrpc": "2.0",
  "method": " get-packet-drop-reason",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "result": {
    "packet": "0010203232..",
    "port": "1",
    "drop-reason": ["vlan-mismatch"],
    "drop-reason-code": "000A2324",
    "drop-location": "dropped-at-pipeline"
  },
  "id": 1
}

```

6.8 API

6.8.1 Overview

This section lists various APIs supported by the Packet Trace feature. A detailed description of each API is available in subsequent sections.

Table 21: Packet Trace APIs

API	Type	Description
configure-packet-trace-feature	Configuration	Configure global options for packet trace and reporting.
get-packet-trace-feature	Status Reporting	Obtain the current configuration for packet trace.
get-packet-trace-lag-resolution	Status Reporting	Request the LAG resolution report for a given packet/packet match criterion.
get-packet-trace-ecmp-resolution	Status Reporting	Request the ECMP resolution report for a given packet/packet match criterion.
get-packet-trace-profile	Status Reporting	Request the complete packet resolution report for a given packet/packet match criterion.
get-packet-drop-reason	Status Reporting	Request the drop profile for a given packet
get-packet-dropctrs	Status Reporting	Request drop statistics and drop profile for live traffic that matches a criterion
cancel-request	Configuration	Stop a previously initiated periodic reporting command.

6.8.2 Configuration/Clear API

6.8.2.1 configure-packet-trace-feature

The `configure-packet-trace-feature` command sets up the Packet Trace functionality on the Agent. The parameters associated with this command are described in [Table 22](#).

Table 22: configure-packet-trace-feature Parameters

Parameter	Type	Description
<code>packet-trace-enable</code>	Boolean	Determines whether the Packet Trace feature should be active or not on the Agent. By default, the Packet Trace feature is inactive on the Agent.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "configure-packet-trace-feature",
  "asic-id": "1",
  "params": {
    "packet-trace-enable": 1
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

6.8.3 Status/Reporting API

6.8.3.1 get-packet-trace-feature

The `get-packet-trace-feature` command is used to retrieve the current configuration of the Packet Trace functionality on the Agent. This API does not require any parameters.

For the response, the Agent returns the same parameters as that of the `configure-packet-trace-feature` command.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-feature",
  "asic-id": "1",
  "params": {},
  "id": 1
}
```

An associated sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-feature",
  "asic-id": "1",
  "version": "1",
  "result": {
    "packet-trace-enable": 1
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

6.8.3.2 get-packet-trace-lag-resolution

The `get-packet-trace-lag-resolution` command is used to retrieve the LAG resolution. There are two sets of parameters to this API depending on the following two possibilities:

- Option 1: The client prefers a trace-profile for an injected debug/PCAP packet.
- Option 2: The client prefers to specify a match criterion and needs a trace-profile for live traffic matching that criterion.⁸

The set of parameters for Option 1 are described in [Table 23](#).

Table 23: get-packet-trace-lag-resolution Option 1 Parameters

Parameter	Type	Description
packet	String	The packet ^a to which the trace-profile is requested. The Requester has to send a full packet to the Agent. The Agent injects this packet into the ingress pipeline as it receives the packet.
port-list	Array of String	Determines the ports on which the trace-profile is requested.
collection-interval	Integer	Determines the periodicity with which the trace-profiles are collected from the ASIC. The units for this parameter are seconds. The default value for this is 60 seconds. The minimum value of interval is 0 second. If the value is 0, the trace-profile report is sent as response to request and periodic reporting is not enabled.
drop-packet	Boolean	Determines whether the ASIC has to drop the packet or forward it once the trace-profile is collected. The default value is true.

a. The packet format is described in [Section 6.3, Packet Format](#).

The set of parameters for Option 2 are described in [Table 24](#).

Table 24: get-packet-trace-lag-resolution Option 2 Parameters

Parameter	Type	Description
src-ip	String	IP address is used to match the source IP address of the packet.
dst-ip	String	IP address is used to match destination IP address of the packet.
protocol	String	Protocol number is used to match the L4 protocol number in the packet.
l4-src-port	String	L4 source port number is used to match the L4 source port of the packet.
l4-dst-port	String	L4 destination port number is used to match the L4 destination port of the packet.
port-list	Array of String	Determines the ports on which trace-profile is requested.
packet-limit	Integer	Determines the maximum number of packet samples for configured interval (as specified by the collection-interval parameter below) to be sent to the collector. Default value is 5. The range is 1 to 10.
collection-interval	Integer	Determines the time period during which maximum of n (as configured using packet-limit) trace-profiles are collected from the ASIC. The units for this parameter are seconds. The default value for this is 10 seconds. The minimum value of interval is 1 second.
drop-packet	Boolean	Determines whether ASIC has to drop the packet or forward once the trace-profile is collected. Default value is true.

8. This method is described in [Section 6.4, Live Traffic Triggered Trace-Profile/Drop-Profile](#).

A sample JSON-RPC request containing the packet is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-lag-resolution",
  "asic-id": "1",
  "params": {
    "packet": "000001000002...",
    "port-list": ["1", "2", "3", "4-10"],
    "collection-interval": 10,
    "drop-packet": 1
  },
  "id": 1
}
```

An associated sample response for the above requests is shown below. The response is asynchronous if the collection interval is non-zero.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-lag-resolution",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "packet-info": [{
    "packet-received-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-ingress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-egress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet": "0010203232.."
  }],
  "report": [
    {
      "port": "1",
      "lag-link-resolution": {
        "lag-id": "2",
        "lag-members": [
          "1",
          "2",
          "3",
          "4"
        ],
        "dst-lag-member": "4"
      }
    },
    {
      "port": "2",
      "lag-link-resolution": {
        "lag-id": "2",
        "lag-members": [
          "1",
          "2",
          "3",
          "4"
        ],
        "dst-lag-member": "4"
      }
    }
  ],
  "id": 1
}
```

A sample JSON-RPC request containing 5-tuple information is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-lag-resolution",
  "asic-id": "1",
  "params": {
    "src-ip": "192.168.1.1",
    "dst-ip": "10.10.1.1",
    "protocol": "4",
    "l4-src-port": "27",
    "l4-dst-port": "35",
    "port-list": ["1", "2", "3", "4-10"],
    "packet-limit": 5,
    "collection-interval": 1,
    "drop-packet": 1
  },
  "id": 1
}
```

The profile information is sent asynchronously for the request if the 5-tuple match criterion is specified.

An associated sample response for the above request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": " get-packet-trace-lag-resolution",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "packet-info": [{
    "packet-received-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-ingress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-egress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet": "0010203232.."
  }],
  "report": [
    {
      "port": "1",
      "lag-link-resolution": {
        "lag-id": "2",
        "lag-members": [
          "1",
          "2",
          "3",
          "4"
        ],
        "dst-lag-member": "4"
      }
    },
    {
      "port": "2",
      "lag-link-resolution": {
        "lag-id": "2",
        "lag-members": [
          "1",
          "2",
          "3",
          "4"
        ],
        "dst-lag-member": "4"
      }
    }
  ],
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

6.8.3.3 get-packet-trace-ecmp-resolution

The get-packet-trace-ecmp-resolution API retrieves the ECMP resolution. There are two sets of parameters to this API depending on the two possibilities below.

- Option 1: The client prefers a trace-profile for a pre-available packet.
- Option 2: The client prefers to specify a match criterion and needs a trace-profile for live traffic matching that criterion.⁹

The set of parameters for Option 1 are described in [Table 25](#).

Table 25: get-packet-trace-ecmp-resolution Option 1 Parameters

Parameter	Type	Description
packet	String	Packet to which trace-profile is requested. The requester has to send a full packet to the Agent, and the Agent will inject this packet into the ingress pipeline as if it received the packet.
port-list	Array of String	Determines the ports on which a trace-profile is requested.
collection-interval	Integer	Determines the periodicity with which the trace-profiles are collected from the ASIC. The units for this parameter are seconds. The default value for this is 60 seconds. The minimum value of interval is 0 second. If the value is 0, the trace-profile report is sent as response to request and periodic reporting is not enabled.
drop-packet	Boolean	Determines whether the ASIC has to drop the packet or forward once the trace-profile is collected. Default value is true.

The set of parameters for Option 2 are described in [Table 26](#).

Table 26: get-packet-trace-ecmp-resolution Option 2 Parameters

Parameter	Type	Description
src-ip	String	IP address is used to match the source IP address of the packet.
dst-ip	String	IP address is used to match destination IP address of the packet.
protocol	String	Protocol number is used to match L4 protocol number in the packet.
l4-src-port	String	Layer-4 source port number is used to match the L4 source port of the packet.
l4-dst-port	String	Layer-4 destination port number is used to match the L4 destination port of the packet.
port-list	Array of String	Determines the ports on which trace-profile is requested.
packet-limit	Integer	Number of samples made by the Agent in the collection interval. Default value is 5. The range is 1 to 10.
collection-interval	Integer	Determines the time period during which maximum of n (as configured using packet-limit) trace-profiles are collected from the ASIC. The units for this parameter are seconds. The default value for this is 10 seconds. The minimum value of interval is 1 second.
drop-packet	Boolean	Determines whether the ASIC has to drop the packet or forward once the trace-profile is collected. Default value is true.

9. This method is described in [Section 6.4, Live Traffic Triggered Trace-Profile/Drop-Profile](#).

A sample JSON-RPC request containing packet is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-ecmp-resolution",
  "asic-id": "1",
  "params": {
    "packet": "000001000002...",
    "port-list": ["1", "2", "3", "4-10"],
    "collection-interval": 10,
    "drop-packet": 1
  },
  "id": 1
}
```

An associated sample response is shown below. The response is asynchronous if the collection interval is non-zero.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-ecmp-resolution",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "packet-info": [{
    "packet-received-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-ingress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-egress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet": "0010203232.."
  }],
  "report": [
    {
      "port": "1",
      "ecmp-link-resolution": [
        {
          "ecmp-group-id": "200256",
          "ecmp-members": [
            ["100005", "3.3.3.2", "15"],
            ["100006", "4.4.4.2", "16"]
          ],
          "ecmp-dst-member": "100005",
          "ecmp-dst-port": "15",
          "ecmp-next-hop-ip": "3.3.3.2"
        }
      ]
    },
    {
      "port": "2",
      "ecmp-link-resolution": [
        {
          "ecmp-group-id": "200256",
          "ecmp-members": [
            ["100005", "3.3.3.2", "15"],
            ["100006", "4.4.4.2", "16"]
          ],
          "ecmp-dst-member": "100005",
          "ecmp-dst-port": "15",
          "ecmp-next-hop-ip": "3.3.3.2"
        }
      ]
    }
  ]
}
```

```

    }
  ],
  "id": 1
}

```

A sample JSON-RPC request containing 5-tuple information is shown below.

```

{
  "jsonrpc": "2.0",
  "method": " get-packet-trace-ecmp-resolution",
  "asic-id": "1",
  "params": {
    "src-ip": "192.168.1.1",
    "dst-ip": "10.10.1.1",
    "protocol": "4",
    "l4-src-port": "27",
    "l4-dst-port": "35",
    "port-list": ["1", "2", "3", "4-10"],
    "packet-limit": 5,
    "collection-interval": 1,
    "drop-packet": 1
  },
  "id": 1
}

```

An associated sample response is shown below. The profile information is sent asynchronously for the request if the 5-tuple match criterion is specified.

```

{
  "jsonrpc": "2.0",
  "method": " get-packet-trace-ecmp-resolution",
  "asic-id": "1",
  "version": "31",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "packet-info": [{
    "packet-received-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-ingress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-egress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet": "0010203232.."
  }],

  "report": [{
    "port": "1",
    "ecmp-link-resolution": [{
      "ecmp-group-id": "200256",
      "ecmp-members": [
        ["100005", "3.3.3.2", "15"],
        ["100006", "4.4.4.2", "16"]
      ]
    },
    "ecmp-dst-member": "100005",
    "ecmp-dst-port": "15",
    "ecmp-next-hop-ip": "3.3.3.2"
  }]
},
  "id": 1
}

```

The Agent acknowledges the command and returns an appropriate error code if needed.

6.8.3.4 get-packet-trace-profile

The get-packet-trace-profile API retrieves the Packet Trace egress resolution. This API obtains the LAG resolution (get-packet-trace-lag-resolution), ECMP resolution (get-packet-trace-ecmp-resolution) and a link resolution from the ASIC.

There are two sets of parameters to this API depending on the following two possibilities:

- Option 1: The client prefers a trace-profile for a pre-available packet.
- Option 2: The client prefers to specify a match criterion and needs a trace-profile for live traffic matching that criterion.¹⁰

The Agent attempts to resolve LAG and ECMP for the packet(s) obtained. The response includes one of the following:

- The LAG resolution (same as response of resolution get-packet-trace-lag-resolution).
- The ECMP resolution (same as response of resolution get-packet-trace-ecmp-resolution).
- The link resolution (egress port information) if the packet could not be traced to either a LAG or ECMP.

The set of parameters for Option 1 are described in [Table 27](#).

Table 27: get-packet-trace-profile Option 1 Parameters

Parameter	Type	Description
packet	String	Packet to which trace-profile is requested. Requester have to send full packet to Agent, Agent will inject this packet into the ingress pipeline as if it receives the packet.
port-list	Array of String	Determines the ports on which trace-profile is requested.
collection-interval	Integer	Determines the periodicity with which the trace-profiles are collected from the ASIC. The units for this parameter are seconds. The default value for this is 60 seconds. The minimum value of interval is 0 second. If the value is 0, the trace-profile report is sent as response to request and periodic reporting is not enabled.
drop-packet	Boolean	Determines whether ASIC has to drop the packet or forward once the trace-profile is collected. Default value is true.

10. This method is described in [Section 6.4, Live Traffic Triggered Trace-Profile/Drop-Profile](#).

The set of parameters for Option 2 are described in [Table 28](#).

Table 28: get-packet-trace-profile Option 2 Parameters

Parameter	Type	Description
src-ip	String	IP address is used to match the source IP address of the packet.
dst-ip	String	IP address is used to match destination IP address of the packet.
protocol	String	Protocol number is used to match layer-4 protocol number in the packet.
l4-src-port	String	Layer-4 source port number is used to match the L4 source port of the packet.
l4-dst-port	String	Layer-4 destination port number is used to match the L4 destination port of the packet.
port-list	Array of String	Determines the ports on which trace-profile is requested.
packet-limit	Integer	Number of samples made by the Agent in the collection interval. Default value is 5. The range is 1 to 10.
collection-interval	Integer	Determines the time period during which maximum of n (as configured using packet-limit) trace-profiles are collected from the ASIC. The units for this parameter are seconds. The default value for this is 10 seconds. The minimum value of interval is 1 second.
drop-packet	Boolean	Determines whether ASIC has to drop the packet or forward once the trace-profile is collected. Default value is true.

A sample JSON-RPC request containing packet is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-profile",
  "asic-id": "1",
  "params": {
    "packet": "000001000002...",
    "port-list": ["1", "2", "3", "4-10"],
    "collection-interval": 10,
    "drop-packet": 1
  },
  "id": 1
}
```

An associated sample response is shown below. The response is asynchronous if the collection interval is non-zero.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-profile",
  "asic-id": "1",
  "version": "3",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "packet-info": [{
    "packet-received-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-ingress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-egress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet": "0010203232.."
  }],
  "report": [
    {
      "port": "1",
      "trace-profile": [
        {

```

```

        "realm": "lag-link-resolution",
        "data": {
            "lag-id": "2",
            "lag-members": ["1", "2", "3", "4"],
            "dst-lag-member": "4"
        }
    },
    {
        "realm": "ecmp-link-resolution",
        "data": [
            {
                "ecmp-group-id": "200256",
                "ecmp-members": [
                    ["100005", "3.3.3.2", "15"],
                    ["100006", "4.4.4.2", "16"]],
                "ecmp-dst-member": "100005",
                "ecmp-dst-port": "15",
                "ecmp-next-hop-ip": "3.3.3.2"
            }
        ]
    },
    {
        "realm": "link-resolution",
        "data": [{
            "dst-port": "15"
        }]
    }
],
{
    "port": "2",
    "trace-profile": [
        {
            "realm": "lag-link-resolution",
            "data": {
                "lag-id": "2",
                "lag-members": ["1", "2", "3", "4"],
                "dst-lag-member": "4"
            }
        },
        {
            "realm": "ecmp-link-resolution",
            "data": [
                {
                    "ecmp-group-id": "200256",
                    "ecmp-members": [
                        ["100005", "3.3.3.2", "15"],
                        ["100006", "4.4.4.2", "16"]],
                    "ecmp-dst-member": "100005",
                    "ecmp-dst-port": "15",
                    "ecmp-next-hop-ip": "3.3.3.2"
                }
            ]
        },
        {
            "realm": "link-resolution",

```

```

        "data": [{
            "dst-port": "15"
        }]
    }
]
},
"id": 1
}

```

A sample JSON-RPC request containing 5-tuple information is shown below.

```

{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-profile",
  "asic-id": "1",
  "params": {
    "src-ip": "192.168.1.1",
    "dst-ip": "10.10.1.1",
    "protocol": "4",
    "l4-src-port": "27",
    "l4-dst-port": "35",
    "port-list": ["1", "2", "3", "4-10"],
    "packet-limit": 5,
    "collection-interval": 1,
    "drop-packet": 1
  },
  "id": 1
}

```

An associated sample response is shown below. The profile information is sent asynchronously for the request if the 5-tuple match criterion is specified.

```

{
  "jsonrpc": "2.0",
  "method": "get-packet-trace-profile",
  "asic-id": "1",
  "version": "13",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "packet-info": [{
    "packet-received-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-ingress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet-egress-time-stamp": "2014-11-18 - 00:15:00 ",
    "packet": "0010203232.."
  }],

  "report": [{
    "port": "1",
    "trace-profile": [{
      "realm": "lag-link-resolution",
      "data": {
        "lag-id": "2",
        "lag-members": ["1", "2", "3", "4"],
        "dst-lag-member": "4"
      }
    }, {
      "realm": "ecmp-link-resolution",
      "data": [{
        "ecmp-group-id": "200256",
        "ecmp-members": [

```

```

        ["100005", "3.3.3.2", "15"],
        ["100006", "4.4.4.2", "16"]],
        "ecmp-dst-member": "100005",
        "ecmp-dst-port": "15",
        "ecmp-next-hop-ip": "3.3.3.2"
    }]
},
{
    "realm": "link-resolution",
    "data": [{
        "dst-port": "15"
    }]
}
}],
    "id": 1
}

```

The Agent acknowledges the command and returns an appropriate error code if needed.

6.8.3.5 get-packet-drop-reason

The get-packet-drop-reason API is used to obtain a drop-profile for the given packet. The set of parameters for this API are described below.

Table 29: get-packet-drop-reason Parameters

Parameter	Type	Description
packet	String	The packet ^a to which the drop-profile is requested. The Requester has to send a full packet to the Agent. The Agent injects this packet into the ingress pipeline.
port	String	Determines the port on which the drop-profile is requested.

a. The packet format is described in [Section 6.3, Packet Format](#).

A sample JSON-RPC request containing the packet is shown below.

```

{
    "jsonrpc": "2.0",
    "method": "get-packet-drop-reason",
    "asic-id": "1",
    "params": {
        "packet": "000001000002...",
        "port": "1"
    },
    "id": 1
}

```

The Agent acknowledges the command and returns an appropriate error code if needed. It returns the drop-profile for the specified packet if there are no errors. The drop-profile is described in [Section 6.7, Deciphering Drop-Profile](#). A Sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": " get-packet-drop-reason",
  "asic-id": "1",
  "version": "6",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "result": {
    "packet": "0010203232..",
    "port": "1",
    "drop-reason": ["vlan-mismatch"],
    "drop-reason-code": "000A2324",
    "drop-location": "dropped-in-pipeline"
  },
  "id": 1
}
```

6.8.3.6 get-packet-dropctrs

The get-packet-dropctrs API is used to track the number of packets getting dropped at different locations in the ASIC, for the ingressing traffic matching a client supplied criterion. Additionally, the API can also be used to obtain drop-profile for some of the matching dropped packets. The set of parameters for this API are described below.

Table 30: get-packet-dropctrs

Parameter	Type	Description
five-tuple	JSON Object	Specifies the flow matching criterion for counting the dropped packets. This is an optional parameter. When not specified, it means all flows must be matched. The JSON object structure is described earlier (in the Live Packet Trace Profile / Drop Profile section) in this document.
port	String	Determines the port on which the drop-counters and/or drop-profile is requested.
duration-in-sec	Integer	Indicates the time period for measuring the number of dropped packets. The units are seconds. The range is 1 – 60 It must be noted that in the event of heavy drops in the ASIC, the counters may overflow when the duration is set to a high value.
send-sample-packets	Boolean	Indicates whether the Agent should send some of the dropped packets to the Client. This is an optional parameter and the default value is 0 (turned off). The sample packets are sent along with the counters, at the expiry of the sampling interval (duration-in-sec).
max-sample-packets	Integer	Valid only when the send-sample-packets is set to 1 (turned on). Indicates the maximum number of sample packets that may be collected for sending to Client. This is an optional parameter and the default value is 25. The range is 1 – 100
include-drop-reasons-for-sample-packets	Boolean	Valid only when the send-sample-packets is set to 1 and max-sample-packets is set to a non-zero value. Indicates that the Agent must send drop-profile along with the dropped-packet samples. This is an optional parameter and the default value is 0. When enabled (set to 1), the Agent will inject the captured packet and gathers the drop-profile.

A sample JSON-RPC request containing the packet is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-packet-dropctrs",
  "asic-id": "1",
  "params": {
    "five-tuple": {
      "src-ip": "192.168.1.1",
      "dst-ip": "10.10.1.1",
      "protocol": "4",
      "l4-src-port": "27",
      "l4-dst-port": "any"
    },
    "port": "10",
    "duration-in-sec": 10,
    "send-sample-packets": 1,
    "max-sample-packets": 5,
    "include-drop-reasons-for-sample-packets": 1
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

The counters and the drop-profile are returned as an asynchronous notification. A sample asynchronous response is shown below. The drop-profile is described in [Section 6.7, Deciphering Drop-Profile](#).

```
{
  "jsonrpc": "2.0",
  "method": " get-packet-dropctrs",
  "asic-id": "1",
  "version": "6",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "result": {
    "counters": {
      "dropped-at-ingress": 1023,
      "dropped-in-pipeline": 212,
      "dropped-at-egress": 345
    },
    "num-sample-packets": 2,
    "packets": [
      {
        "packet": "0010203232..",
        "port": "10",
        "drop-reason": ["vlan-mismatch"],
        "drop-reason-code": "000A2324",
        "drop-location": "dropped-in-pipeline"
      },
      {
        "packet": "0020203232..",
        "port": "10",
        "drop-reason": ["vlan-mismatch"],
        "drop-reason-code": "000A2324",
        "drop-location": "dropped-in-pipeline"
      }
    ]
  },
  "id": 1
}
```

Chapter 7: Black Hole Detection

7.1 Overview¹¹

Black holes refer to logical places in the network where incoming or outgoing traffic is silently discarded. A black hole condition arises when the traffic is directed towards an incorrect path, which causes the packets to traverse the same set of nodes in a loop. This continues until the packet's Time-to-Live (TTL) expires, at which point the packet is discarded. This condition is called *black holing* as the packet fails to reach the destination node.

A black hole in the data plane is a symptom of an inconsistent configuration of one or more routing tables in the network data plane. Such inconsistency may be the result of a benign condition that will heal by itself, such as a link or switch failure, and the awareness of the failure has not fully propagated to the entire network. The inconsistency might also be a persistent condition such as a hardware bug, a software defect, or a configuration mistake.

The BroadView Agent allows clients to monitor for potential black holes in the network. Upon appropriate configuration, the Agent notifies the client of a likely black hole, while also providing the specific packets that are being black holed. It allows network operators to separate the benign from the problematic causes, take immediate corrective action to avoid persistent packet loss, and help isolate the root cause.

Optionally, the BroadView Agent can also send samples of black holed packets to an sFlow receiver.

7.2 Black Hole Ports

In a typical spine-leaf topology, the uplink ports on the leaf switches connecting back to spine switches are called *spine ports*. Packets that enter leaf switches from spine switches tend not to go back to spine switches. If a packet comes into a leaf switch from one of the spine ports and egresses from another of the spine ports, it is usually a symptom of black holing. Such packets are likely to be forward within the same spine-leaf-spine loop and will eventually be discarded.

In the context of this feature, the set of ports where the packet cannot both ingress and egress are called *black hole ports*. Typically, these are the spine ports.

11. A white paper describing the Black Hole Detection and Avoidance feature is available on the Broadcom.com website. Refer to *Black Hole Detection by BroadView™ Instrumentation Software Abstract* (BHD-WP101-R.pdf).

7.3 Reporting Black Hole Events

The Broadview Instrumentation Agent monitors the black hole ports and reports each black holing event to the configured client. The event report includes the following information:

- Ingress port
- Egress port
- Packet that caused the event

This information enables the client to take appropriate corrective action. It is also possible to configure the Agent to send the black holing packets to an sFlow receiver connected to one of the switch ports.

7.4 Configurability

The following configurability is provided by the BHD feature.

- The BHD feature can be enabled or disabled on the Agent.
- The client provides the list of black hole ports.
- The BHD feature can be configured to monitor for black holing packets and send black holing event reports to the client (as notification reports).
- The client can configure a watermark level for black holing packets (in packets/sec) to ignore transient conditions.
- The BHD feature can be configured to send the black holing packets to an sFlow receiver connected to one of the switch ports.

Some of the BHD capabilities, such as reporting to an sFlow receiver, require specific ASIC capabilities. If the client configuration cannot be supported on the underlying hardware, the Agent returns an appropriate error.

7.5 API

7.5.1 Overview

This section lists APIs supported by the Black Hole Detection (BHD) feature. A detailed description of each of the APIs is available in subsequent sections.

Table 31: Black Hole Detection APIs

API	Type	Description
black-hole-detection-enable	Configuration	Configure global options for BHD.
configure-black-hole	Configuration	Defines a port set constituting a black hole and reporting parameters.
cancel-black-hole	Configuration	Clears a previously set black hole configuration and reporting parameters.
get-black-hole-detection-enable	Status Reporting	Obtain current BHD configuration.
get-black-hole	Status Reporting	Obtain the current black hole configuration and reporting parameters.
get-black-hole-event-report	Notification	Report from Agent indicating a black-holed packet event.
get-sflow-sampling-status	Status Reporting	Obtain the sFlow reporting status.
cancel-request	Configuration	Stop a previously initiated periodic reporting command.

7.5.2 Configuration API

7.5.2.1 black-hole-detection-enable

The `black-hole-detection-enable` command enables the Black Hole Detection functionality on the switch. The parameter associated with this command are described in [Table 32](#).

Table 32: black-hole-detection-enable Parameter

Parameter	Type	Description
<code>enable</code>	Boolean	When enabled, the BroadView Agent enables the Black Hole Detection feature. When disabled, Black Hole Detection feature is disabled on the Agent and any existing configuration is discarded. By default, the feature is turned OFF.

A sample JSON-RPC request to enable BHD is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "black-hole-detection-enable",
  "asic-id": "1",
  "params": {
    "enable": 1
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

7.5.2.2 configure-black-hole

The `configure-black-hole` command configures the Black Hole Detection functionality on the switch. The parameters associated with this command are described in [Table 33](#).

Table 33: configure-black-hole Parameters

Parameter	Type	Description
<code>port-list</code>	String	List of ports where the traffic is monitored for black holes.
<code>sampling-method</code>	String Enumeration	Sampling method used to sample packets. Supported values are <i>agent</i> and <i>sflow</i> .

Sampling parameters for Agent sampling are provided in [Table 34](#).

Table 34: configure-black-hole Agent Sampling Parameters

Parameter	Type	Description
<code>water-mark</code>	Integer	This represents the traffic rate in packet/second, above which traffic is considered as black holed. Sampling starts only after the watermark level is crossed. The default is 1024.
<code>sample-periodicity</code>	Integer	Time interval in seconds. The default is 10 seconds.
<code>sample-count</code>	Integer	Number of samples to be sent for every <i>sample-periodicity</i> interval. The default is 2. If sample-count is set to 0, then sampling is stopped.

Sampling parameters for sFlow sampling are provided in [Table 35](#).

Table 35: configure-black-hole sFlow Sampling Parameters

Parameter	Type	Description
<code>mirror-port</code>	String	Port number on which sFlow encapsulated sample packet is sent out.
<code>vlan-id</code>	Integer	The VLAN identifier of the sFlow encapsulation header.
<code>destination-ip</code>	String	Destination IP address in the sFlow encapsulation header.
<code>source-udp-port</code>	Integer	Source UDP port number in the sFlow encapsulation header.
<code>destination-udp-port</code>	Integer	Destination UDP port number in the sFlow encapsulation header.
<code>sample-pool-size</code>	Integer	Represents the packet pool size for sampling. One packet is sampled out of <i>sample-pool-size</i> packets. Minimum is 1024. The default is 1024.

A sample JSON-RPC request to configure BHD using Agent sampling is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "configure-black-hole ",
  "asic-id": "1",
  "params": {
    "port-list": ["1", "2", "3", "4"],
    "sampling-method": "agent",
    "sampling-params": {
      "water-mark": 200,
      "sample-periodicity": 15,
      "sample-count": 10
    }
  },
  "id": 1
}
```

A sample JSON-RPC request to configure BHD using sFlow agent sampling is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "configure-black-hole ",
  "asic-id": "1",
  "params": {
    "port-list": ["1", "2", "3", "4"],
    "sampling-method": "sflow",
    "sampling-params": {
      "encapsulation-params": {
        "vlan-id": 1,
        "destination-ip": "1.1.1.1",
        "source-udp-port": 1234,
        "destination-udp-port": 4321
      },
      "mirror-port": "10",
      "sample-pool-size": 1024
    }
  },
  "id": 1
}
```

The Agent acknowledges the command, and returns an appropriate error code if needed.

7.5.2.3 cancel-black-hole

The `cancel-black-hole` API removes all black hole configurations on the switch. There are no parameters associated with this command.

A sample JSON-RPC request to cancel black hole is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "cancel-black-hole",
  "asic-id": "1",
  "params": {
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

7.5.3 Status/Reporting API

7.5.3.1 get-black-hole-detection-enable

The `get-black-hole-detection-enable` API retrieves the current status of the BHD feature on the Agent. There are no parameters to this API. The parameters for response are same as that of [black-hole-detection-enable](#) API.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-black-hole-detection-enable",
  "asic-id": "1",
  "params": {},
  "id": 1
}
```

An associated sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-black-hole-detection-enable",
  "asic-id": "1",
  "version": "2",
  "result": {
    "enable": 1
  },
  "id": 1
}
```

7.5.3.2 get-black-hole

The `get-black-hole` API retrieves the current configuration of the Black Hole Detection feature on the Agent. There are no parameters to this API. The parameters for response are same as that of [configure-black-hole](#) API.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-black-hole",
  "asic-id": "1",
  "params": {},
  "id": 1
}
```

An associated sample response is shown below for Agent sampling.

```
{
  "jsonrpc": "2.0",
  "method": "get-black-hole",
  "asic-id": "1",
  "version": "2",
  "result": {
    "port-list": ["1", "2", "3", "4"],
    "sampling-method": "agent",
    "sampling-params": {
      "water-mark": 200,
      "sample-periodicity": 15,
      "sample-count": 10
    }
  },
  "id": 1
}
```

An associated sample response is shown below for SFlow sampling.

```
{
  "jsonrpc": "2.0",
  "method": "get-black-hole",
  "asic-id": "1",
  "version": "2",
  "result": {
    "port-list": ["1", "2", "3", "4"],
    "sampling-method": "sflow",
    "sampling-params": {
      "encapsulation-params": {
        "vlan-id": 1,
        "destination-ip": "1.1.1.1",
        "source-udp-port": 1234,
        "destination-udp-port": 4321
      },
      "mirror-port": "10",
      "sample-pool-size": 1
    }
  },
  "id": 1
}
```

7.5.3.3 get-sflow-sampling-status

The get-sflow-sampling-status API retrieves the current sFlow status for black-holed traffic. The parameter associated with this command is described in [Table 36](#).

Table 36: get-sflow-sampling-status Parameter

Parameter	Type	Description
port-list	String	List of ports on which the sFlow sampling status is requested.

The REST JSON request message for the same is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-sflow-sampling-status",
  "asic-id": "1",
  "params": {
    "port-list": ["1", "2"]
  },
  "id": 1
}
```

The parameters associated with the response are described in [Table 37](#).

Table 37: get-sflow-sampling-status Response Parameters

Parameter	Type	Description
port	String	Port number.
sflow-sampling-enabled	Boolean	sFlow sampling status.
sampld-packet-count	Integer	Total number of packets sampled since sFlow sampling is enabled.
black-holed-packet-count	Integer	Total number of packets black-holed since sFlow sampling is enabled.

An associated sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-sflow-sampling-status",
  "asic-id": "1",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "version": "2",
  "report": {
    "data": [{
      "port": "2",
      "sflow-sampling-enabled": 1,
      "sampld-packet-count": 100,
      "black-holed-packet-count": 1000
    }, {
      "port": "3",
      "sflow-sampling-enabled": 1,
      "sampld-packet-count": 200,
      "black-holed-packet-count": 2000
    }
  ],
  "id": 1
}
```

7.5.4 Notifications

7.5.4.1 get-black-hole-event-report

The REST JSON message `get-black-hole-event-report` is used to send asynchronous black hole event reports, including the sampled packet, to the client. This is applicable only when the Agent is configured in Agent sampling mode.

The parameters associated with this notification are described in [Table 38](#).

Table 38: get-black-hole-event-report Parameters

Parameter	Type	Description
<code>ingress-port</code>	String	Port number on which the packet has entered the switch.
<code>egress-port-list</code>	String	List of egress ports where the packet would be sent out.
<code>black-holed-packet-count</code>	Integer	Total number of packets black holed.
<code>sample-packet</code>	String	Sampled packet in PCAP format with BASE 64 encoding.

A sample notification message is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-black-hole-event-report",
  "asic-id": "1",
  "version": "2",
  "time-stamp": "2014-11-18 - 00:15:04 ",
  "report": {
    "ingress-port": "1",
    "egress-port-list": ["2", "3"],
    "black-holed-packet-count": 100,
    "sample-packet": "0010203232.."
  }
}
```

Sampled packets larger than 1588 bytes are truncated and might appear as malformed at the client side.

Chapter 8: Flow Tracker

8.1 Overview

The BroadView Flow Tracker feature works in conjunction with the Flow Tracker Embedded Application to provide deep visibility into various flows passing through the switch.

A *flow* is identified by the standard five-tuple match criterion: the pair of source and destination IP addresses, the pair of source and destination ports, and the protocol field (in the IP header).

A *flow group* identifies a set of flows. The following types of flow groups are supported.

- **Five-tuple-based flow group.** This type of flow group specifies the standard five-tuple (for a flow), but one or more fields are specified as *wildcards*. Multiple flows from the live traffic may belong to this flow-group type.
- **Egress-port based flow group.** All flows destined to pass through the specified egress ports belong to this flow group. Up to eight egress ports may be specified.
- **Ingress-port based flow group.** All flows coming into the switch through the specified ingress ports belong to this flow group. Up to eight ingress ports may be specified.
- **Congestion-based flow group.** All flows destined to egress from the ports that are experiencing the most congestion in the switch. The top eight egress ports experiencing congestion are automatically included for monitoring. Congestion is measured in two parameters: number of dropped packets and buffer utilization.

Clients can specify the flow-groups for monitoring using the REST API provided by the BroadView Flow Tracker feature. The switch monitors all flows belonging to the specified flow groups and sends the collected instrumentation data as IPFIX records to an IPFIX collector entity.

The rest of this chapter describes the various APIs that are supported as part of this feature.

8.2 Communication Exchanges

8.2.1 Information

The switch sends the following types information to a receiving IPFIX entity (Collector):

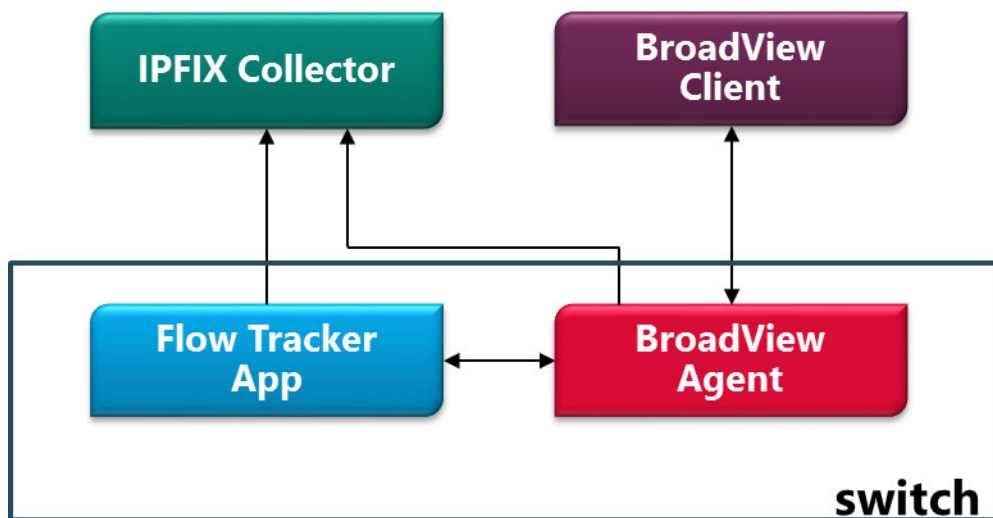
- **IPFIX Record Messages.** The IPFIX records are sent using the UDP transport. Two types of record messages are sent by the switch.
 - **Flow Group Records:** The switch sends the flow-group record messages to the Collector at a periodic interval (configurable). These records contain the following data, recorded in the interval:
 - Flow Group Information
 - Number of Packets and Octets
 - Number of TCP Connections made/terminated.
 - **Flow Records:** The switch sends the flow record messages for all the active flows to the Collector at a periodic interval (configurable). These records contain the following data, recorded in the interval.
 - Observation timestamp
 - For each of the flows:
 - Flow Information
 - Number of packets and octets

- **IPFIX Record Data Set Templates.** The IPFIX record templates are sent using the UDP transport. The following templates (for the record data sets) are sent by the switch to the Collector. These templates are sent every time a flow group is configured on the switch and periodically thereafter (configurable).
 - Flow Group Record Data Set Template
 - Flow Record Data Set Template

8.2.2 Entities

There are four entities involved in an end-to-end feature realization, as [Figure 2](#) shows.

Figure 2: End-to-End Feature Realization Entities



The entities include:

- **IPFIX collector** receives IPFIX templates and the IPFIX messages. The IPFIX collector is responsible for processing. It is responsible for providing the necessary user interface for monitoring the flow data being received from the switch. It may be possible that more than one IPFIX collector is present. In such a configuration, the switch sends a copy of the IPFIX record to each configured Collector.
- **BroadView Client** configures the BroadView Agent via the Agent provided REST API. It is responsible for setting up the Flow Tracker feature on the switch (via the Agent), including the Collector credentials (such as IP address, port number, and so on).
- **Flow Tracker App** runs on the embedded CPU on the switch. It is responsible for sending the IPFIX flow record messages and associated templates periodically to the IPFIX collector.
- **BroadView Agent** runs on the host CPU of the switch. It has multiple communication responsibilities as follows:
 - It configures the Flow Tracker application on the embedded CPU.
 - It sends the IPFIX flow group record messages and associated templates periodically to the IPFIX collector.

8.3 IPFIX Templates

8.3.1 Flow-Group Record Data Set Template

The flow-group record messages are sent by the BroadView Agent, at periodic intervals, to the configured IPFIX collector(s). The IPFIX template¹² that describes these flow-group record data set is provided below¹³.

If any of the fields are not applicable, or wildcards (match-any, ignored) for a flow group, they must not be included in the template record and the group data records when sent to Collector. For example, a group record for a 5-tuple flow-group does not include the field `ingressPhysicalInterface`. Also, the group data/template records for a 5-tuple flow-group, which does not specify source IP address (match any), does not include the `sourceIPv4Prefix` field. The `Field Count` in the template record must be set appropriately.

Each flow-group has a unique flow-group record template based on the configuration. The template IDs for the flow group records start from 301.

SetId = 2	Length = 56
Template Id = 301	Field Count = 11
groupId (Enterprise Element) = 501	Field Length = 1
Enterprise Number	
sourceIPv4Prefix = 44	Field Length = 4
destinationIPv4Prefix = 45	Field Length = 4
sourceTransportPort = 7	Field Length = 2
destinationTransportPort = 11	Field Length = 2
protocolIdentifier = 4	Field Length = 1
ingressPhysicalInterface = 252	Field Length = 4
egressPhysicalInterface = 253	Field Length = 4
originalFlowsPresent = 375	Field Length = 8
packetTotalCount = 86	Field Length = 8
octetTotalCount = 85	Field Length = 8
tcpSynTotalCount = 218	Field Length = 8
tcpFinTotalCount = 219	Field Length = 8
tcpRstTotalCount = 220	Field Length = 8

12. The field references are as per the IANA assignments.

13. The templates in this document follow the RFC style where each row corresponds to a 32-bit word, divided into two 16-bit halves.

8.3.2 Flow Record Data Set Template

The flow record data sets are aggregated and sent by the Flow Tracker App at periodic intervals to the configured IPFIX collector(s). The IPFIX template that describes these flow record data set is provided below.

SetId = 2	Length = 52
Template Id = 257	Field Count = 10
groupId (Enterprise Element) = 501	Field Length = 1
Enterprise Number	
flowStartMilliseconds = 152	Field Length = 8
observationTimeMilliseconds = 323	Field Length = 8
sourceIPv4Address = 8	Field Length = 4
destinationIPv4Address = 12	Field Length = 4
sourceTransportPort = 7	Field Length = 2
destinationTransportPort = 11	Field Length = 2
protocolIdentifier = 4	Field Length = 1
packetTotalCount = 86	Field Length = 4
octetTotalCount = 85	Field Length = 6

8.4 Configurability

The following configurability is provided by the Flow Tracker feature. It may be noted that this configurability/status retrieval is via the REST API. Flow instrumentation data sent as IPFIX records is not included in this list.

- The Flow Tracker feature can be enabled/disabled on the switch.
- The Flow Tracker feature can provide the current capabilities of the switch, in terms of flow-related instrumentation possibilities and current and maximum values for various scalability parameters.
- One or more IPFIX collectors can be configured on the switch. Likewise, a previously configured IPFIX collector can be removed.
- New flow-groups can be added for monitoring. Flow-groups added previously may be removed.
- The Flow Tracker feature can provide statistics for currently active flow-groups.

8.5 API

8.5.1 Overview

This section lists various API supported by the Flow Tracker feature. A detailed description of each of the API is available in subsequent sections.

Table 39: Flow Tracker APIs

API	Type	Description
configure-ft-feature	Configuration	Configure Flow Tracker
create-ft-collector	Configuration	Add an IPFIX collector
remove-ft-collector	Configuration	Remove an IPFIX collector
create-ft-flowgroup	Configuration	Add a new flow-group for monitoring
remove-ft-flowgroup	Configuration	Remove the specified flow-group from monitoring
clear-ft-flowgroup-statistics	Configuration	Clear the statistics for one or all flow-groups
get-ft-collector	Status / Reporting	Retrieve the configuration for one or all collectors
get-ft-flowgroup	Status / Reporting	Retrieve the configuration for one or all flow-groups
get-ft-flowgroup-statistics	Status / Reporting	Retrieve the statistics for one or all flow-groups
get-ft-capabilities	Status / Reporting	Retrieve the Flow Tracker Capabilities and Scaling Parameters
get-ft-status	Status / Reporting	Retrieve the Flow Tracker Status
cancel-request^a	Configuration	Cancel a previously initiated periodic report request

a. This API is a generic BroadView API and is used to cancel a previously invoked API. It is not described in this feature specification

8.5.2 Configuration API

8.5.2.1 configure-ft-feature

The `configure-ft-feature` API sets up the Flow Tracker functionality on the Agent. The following table describes the parameters associated with this API.

Table 40: configure-ft-feature API Parameters

Parameter	Type	Description
<code>ft-enable</code>	Boolean	Determines whether FT feature should be active on the agent. If the feature is disabled from a previously enabled state, all Flow Tracker configuration is removed from the switch. By default, the Flow Tracker feature is inactive on the agent.
<code>aging-time</code>	Integer	Optional Parameter. Indicates the aging time for the flows being tracked. Once the aging time expires, all the counters for the flow are purged. Note that there may be some slight deviation in the actual aging value from the configured setting, depending on the implementation. The units are in seconds. Default value is 60 (1min) Range is 60 to 3600 (one hour)
<code>ipfix</code>	Object	The IPFIX-related parameters are grouped into this JSON object. The parameters within this object are listed in Table 41 . This is an optional parameter object. If not supplied at the first configuration, default values are assumed. If not supplied at a later configuration, any existing values for the parameters are retained.

IPFIX parameters provided via the `ipfix` JSON object are listed in the following table.

Table 41: IPFIX Parameters

Parameter	Type	Description
<code>template-update-periodicity</code>	Integer	Optional parameter. Indicates the periodicity with which the IPFIX templates are sent to the collectors. The units are in seconds. Default value is 900 (15min) Range is 1 to 3600 (one hour)
<code>flow-records-src-port</code>	Integer	Optional parameter. Indicates the source port used for reporting Flow Records to the collector. Default value is 8081
<code>flow-group-records-src-port</code>	Integer	Optional parameter. Indicates the source port used for reporting flow group Records to the collector. Default value is 8082
<code>flow-record-observation-domain</code>	Integer	Optional parameter. Indicates the observation domain used while reporting Flow Records to the collector. Default value is 1
<code>flow-group-record-observation-domain</code>	Integer	Optional parameter. Indicates the observation domain used while reporting flow group records to the collector. Default value is 2

Table 41: IPFIX Parameters (Continued)

Parameter	Type	Description
flow-records-update-periodicity	Integer	Optional parameter. Indicates the frequency (in milliseconds) with which the counters must be sampled and reported to the collector. Default value is 100 (100 ms) Range is 100 to 100000 (100 ms to 1 sec)
flowgroup-records-update-periodicity	Integer	Optional parameter. Indicates the frequency (in seconds) with which the counters must be sampled and reported to the collector. Default value is 5 (5s) Range is 1 to 180 (1s to 3 min)

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "configure-ft-feature",
  "asic-id": "1",
  "params": {
    "ft-enable": 1,
    "ipfix": {
      "template-update-periodicity": 600,
      "flow-records-src-port": 8085,
      "flow-group-records-src-port": 8086
    }
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

8.5.2.2 create-ft-collector

The `create-ft-collector` API adds a collector for receiving the Flow Instrumentation records.

This API merely adds the knowledge of the collector to the Agent. The collector must be associated with a flow-group for receiving the Flow Instrumentation records for the flow-group. UDP is the only supported transport protocol used for exchanging flow instrumentation records¹⁴.

The following table describes the parameters associated with this API.

Table 42: create-ft-collector API Parameters

Parameter	Type	Description
name	String	A string assigned as the name for the collector. Future interactions, such as associating the collector with a specific flow-group will use the <code>name</code> to reference the collector.
protocol	String Enumeration	Protocol that the collector understands for receiving the Instrumentation Records. Valid entries are shown below. Only one protocol is supported currently. IPFIXv10

14. Some implementations may not support multiple collectors to be configured. In such cases, an attempt to create more than one collector returns error. The response to the `get-ft-capabilities` API indicates whether an implementation supports multiple collectors.

Table 42: create-ft-collector API Parameters (Continued)

Parameter	Type	Description
ipaddress	IP Address	IP address of the collector, in a string format. Only IPv4 addresses are supported. It is assumed that the switch has the necessary information to reach the collector.
port	Integer	Optional parameter indicating the UDP port on which the Instrumentation Records are to be sent. If not specified, the port number 2055 is used.
vlan-tag	Integer	Optional parameter. When specified, this value will be added as a VLAN tag to the record packets being sent out. By default, untagged packets are sent.
priority	Integer	Optional parameter indicating the specific priority with which the record packets will be processed by the underlying silicon while transmitting. Usually this number gets mapped to a specific CoS-queue. The default priority is 5.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "create-ft-collector",
  "asic-id": "1",
  "params": {
    "name": "ntop-srv1",
    "protocol": "IPFIXv10",
    "ipaddress": "10.2.2.4",
    "port": 4740
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

8.5.2.3 remove-ft-collector

The `remove-ft-collector` API removes a collector. For a collector to be removed, it must not be currently associated with a flow-group being monitored.

The following table describes the parameters associated with this API.

Table 43: remove-ft-collector API Parameters

Parameter	Type	Description
name	String	Name of the collector to be removed. This collector must have been created previously via the <code>create-ft-collector</code> API. The special keyword <code>all</code> can be used to remove all collectors created thus far.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "remove-ft-collector",
  "asic-id": "1",
  "params": {
    "name": "ntop-srv1"
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

8.5.2.4 create-ft-flowgroup

The `create-ft-flowgroup` API adds a flow-group for monitoring in the switch and associates with one or more collectors for receiving the instrumentation records.

Once a flow-group is created the following events take place:

- The Agent sends the IPFIX templates to the collector(s).
- The switch starts monitoring flows belonging to the flow-group and starts reporting instrumentation records to the collectors.

The following table describes the parameters associated with this API.

Table 44: create-ft-flowgroup API Parameters

Parameter	Type	Description
name	String	A string assigned as the name for the Flow-Group. Future interactions will use the name to reference the flow-group.
id	Integer	A unique identifier used to reference the group in the IPFIX records sent to the collector. Range is 1 to 255

Table 44: create-ft-flowgroup API Parameters (Continued)

Parameter	Type	Description
type	String Enumeration	Type of the flow group ^a . One of the following strings must be specified for this parameter: <ul style="list-style-type: none"> ■ five-tuple ■ egress-ports ■ ingress-ports ■ congestion
collector-list	Array of Strings	Provides a list of the collectors, identified by their names. The switch sends the flow instrumentation records for this flow-group to all the collectors in this list.
load-balancing-mode	String Enumeration	Optional parameter. When multiple collectors are associated, this parameter determines how the Flow Instrumentation Records are load-balanced across them. Valid entries are shown below. Only one protocol is currently supported. DUPLICATE If this parameter is not specified, the switch duplicates the flow records for each collector.
actions	Object	This JSON object provides all the actions to be taken for the matching flows (Described below).
five-tuple	Object	This JSON object must be specified if the type parameter (above) is specified as five-tuple. Otherwise this field is ignored. This JSON object is created by combining the five-tuple fields that constitute the flow-group.
egress-ports	Object	This JSON object must be specified if the type parameter (above) is specified as egress-ports. Otherwise this field is ignored. This JSON object provides the flow match criterion for the egress port-based flow groups.
ingress-ports	Object	This JSON object must be specified if the type parameter (above) is specified as ingress-ports. Otherwise this field is ignored. This JSON object provides the flow match criterion for the Ingress port-based flow groups.
congestion	Object	This JSON object must be specified if the type parameter (above) is specified as congestion. Otherwise this field is ignored. This JSON object provides the flow match criterion for the congestion-based flow groups.

- a. A given implementation may not support all types of flow groups. The response to the get-ft-capabilities API indicates the types of flow groups that the implementation supports.

The following table describes the parameters that constitute the actions JSON object (that define all actions to be taken for the flows matching the flow group).

Table 45: actions JSON Object Parameters

Parameter	Type	Description
report-flow-records	Boolean	Optional parameter. This parameter indicates whether per-flow records must be sent to the IPFIX collector. Default value is true. (Flow records sent to the collector).
report-group-records	Boolean	Optional parameter. In addition to the per-flow records, the BroadView Agent can also send aggregated data belonging to a flow-group. This parameter indicates whether group level records must be sent to the IPFIX collector. Default value is true. (Group level records are sent to the collector).
report-sample-flows	Boolean	Optional parameter. If specified as true, up to num-sample-flows number of packets sampled from the flows matching the flow-group are sent to the BroadView client. Default value is false. (No sampled flows are sent to the BroadView client).
num-sample-flows	Integer	Optional parameter. Valid only when report-sample-flows is true. Default value is 100 (100 sampled flows are sent to the BroadView client). Range is 1 to 1000.
sample-interval	Integer	Optional parameter. Valid only when report-sample-flows is true. Units are in seconds. Default value is 1 (num-sample-flows flows are sampled in 1 sec duration and sent to the BroadView client). Range is 1 to 100. If sample-interval is expired and num-sample-flows flows could not be sampled, only as many as captured is sent to the BroadView client and sampling session is closed.

The five-tuple fields that constitute the flow-group are provided in [Table 46](#). The client must specify and provide valid values for at least two fields among these five. The fields that are not specified are considered as *wildcards*. A wildcard for a field can also be specified by providing the value as the string any.

Table 46: Flow Group Five-Tuple Fields

Parameter	Type	Description
src-ip	String	IP address (along with mask) is used to match the source IP address of the packet.
dst-ip	String	IP address (along with mask) is used to match destination IP address of the packet.
protocol	String	Protocol number is used to match the L4 protocol number in the packet.
l4-src-port	String	L4 source port number is used to match the L4 source port of the packet
l4-dst-port	String	L4 destination port number is used to match the L4 destination port of the packet

The following table describes the parameters that constitute the egress-ports JSON object for creating the egress-ports-based flow group.

Table 47: egress-ports JSON Object Parameters

Parameter	Type	Description
port-list	Array of Strings	An array of all egress ports to be included for monitoring flows. Up to 8 ports may be specified.

The following table describes the parameters that constitute the ingress-ports JSON object for creating the ingress-ports-based flow group.

Table 48: ingress-ports JSON Object Parameters

Parameter	Type	Description
port-list	Array of Strings	An array of all Ingress ports to be included for monitoring flows. Up to 8 ports may be specified.

The following table describes the parameters that constitute the congestion JSON object for creating the congestion-based flow group.

Table 49: congestion JSON Object Parameters

Parameter	Type	Description
num-ports	Integer	Maximum number of ports experiencing congestion to be considered for monitoring. Up to eight ports can be monitored.
min-buffer-utilization	Integer	Minimum buffer utilization, expressed as a percentage of allocated buffer, to be considered as congestion.

A sample JSON-RPC request that creates a five-tuple-based flow group to gather flow instrumentation records for all incoming HTTP connections is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "create-ft-flowgroup",
  "asic-id": "1",
  "params": {
    "name": "httpsrv1",
    "type": "five-tuple",
    "id": 456,
    "collector-list": ["ntop-srv1"],
    "five-tuple" : {
      "dst-ip": "10.2.2.4/32",
      "l4-dst-port": "80"
    }
  },
  "id": 1
}
```

A sample JSON-RPC request that creates an egress ports based flow group to gather flow instrumentation records for the traffic egressing the spine links is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "create-ft-flowgroup",
  "asic-id": "1",
  "params": {
    "name": "httpsrv1",
    "id": 456,
    "type": "egress-ports",
    "collector-list": ["ntop-srv1"],
    "egress-ports" : {
      "port-list": ["1", "2", "3", "4"]
    }
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

8.5.2.5 remove-ft-flowgroup

The `remove-ft-flowgroup` API removes a previously configured flow group. Once a flow group is removed, monitoring will be stopped for any flows belonging to this flow group. Note that the IPFIX collector may receive some transient flow/flow-group records after the API is invoked.

The following table describes the parameters associated with this API.

Table 50: remove-ft-flowgroup API Parameters

Parameter	Type	Description
name	String	Name of the flow group to be removed. This flow group must have been created previously via the <code>create-ft-flow</code> group API. The special keyword <code>all</code> can be used to remove all flow groups created thus far.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "remove-ft-flowgroup",
  "asic-id": "1",
  "params": {
    "name": "httpsrv1"
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

8.5.2.6 clear-ft-flowgroup-statistics

The `clear-ft-flowgroup-statistics` API clears all counters for a flow group. If the statistics are cleared, the IPFIX records being sent to the collector also reflect this event (clearing).

The following table describes the parameters associated with this API.

Table 51: clear-ft-flowgroup-statistics API Parameters

Parameter	Type	Description
name	String	Name of the flow group to clear the statistics for. This flow group must have been created previously via the <code>create-ft-flowgroup</code> API. The special keyword <code>all</code> can be used to clear the statistics for all the flow groups created thus far.
clear-group-records	Boolean	Indicates whether the group level records need to be cleared. Optional parameter, if not specified, the default value of true is assumed.
clear-flow-records	Boolean	Indicates whether the flow level records need to be cleared. Optional parameter, if not specified, the default value of true is assumed.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "clear-ft-flowgroup-statistics",
  "asic-id": "1",
  "params": {
    "name": "httpsrv1"
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed.

8.5.3 Status/Reporting API

8.5.3.1 get-ft-collector

The `get-ft-collector` API retrieves the configuration data associated with a specified collector.

The following table describes the parameters associated with this API.

Table 52: get-ft-collector API Parameters

Parameter	Type	Description
name	String	Name of the collector to retrieve the data for. This collector must have been created previously via the <code>create-ft-collector</code> API. The special keyword <code>all</code> can be used to retrieve the data for all the collectors created thus far.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-collector",
  "asic-id": "1",
  "params": {
    "name": "ntop-srv1"
  },
  "id": 1
}
```

The JSON response includes all parameters listed in the `create-ft-collector` API, as applicable for the collector. The parameter `vlan-tag` is included in the response only when the flow/flow group records are configured to be VLAN tagged.

A sample JSON response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-collector",
  "asic-id": "1",
  "result": [{
    "name": "ntop-srv1",
    "protocol": "IPFIXv10",
    "ipaddress": "10.2.2.4",
    "port": 4740,
    "priority": 5,
  }],
  "id": 1
}
```

8.5.3.2 get-ft-flowgroup

The `get-ft-flowgroup` API retrieves the configuration data associated with a specified flow group.

The following table describes the parameters associated with this API.

Table 53: get-ft-flowgroup API Parameters

Parameter	Type	Description
name	String	Name of the flow group to retrieve the data for. This flow group must have been created previously via the <code>create-ft-flowgroup</code> API. The special keyword <code>all</code> can be used to retrieve the data for all the flow groups created thus far.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-flowgroup",
  "asic-id": "1",
  "params": {
    "name": "httpsrv1"
  },
  "id": 1
}
```

The JSON response includes all parameters listed in the `create-ft-flowgroup` API, as applicable for the flow group.

A sample JSON response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-flowgroup",
  "asic-id": "1",
  "version": "4",
  "result": [{
    "name": "httpsrv1",
    "id": 456,
    "type": "egress-ports",
    "collector-list": ["ntop-srv1"],
    "egress-ports": {
      "port-list": ["1", "2", "3", "4"]
    }
  }],
  "id": 1
}
```

8.5.3.3 get-ft-flowgroup-statistics

The `get-ft-flowgroup-statistics` API retrieves statistics for a flow group.

The following table describes the parameters associated with this API.

Table 54: get-ft-flowgroup-statistics API Parameters

Parameter	Type	Description
<code>name</code>	String	Name of the flow group to retrieve the statistics for. This flow group must have been created previously via the <code>create-ft-flowgroup</code> API. The special keyword <code>all</code> can be used to retrieve the statistics for all the flow groups created thus far.
<code>report-periodicity</code>	Integer	Optional parameter. When specified and has a non-zero value, the statistics are reported at a periodicity provided by this parameter. Units are in seconds. Range is 1 to 3600. By default, periodic reporting is turned off. If this parameter is not specified, the statistics are provided as an immediate response to the API.

The response is an array of statistics objects, which have the parameters [Table 55](#) describes. The statistics are at the group level, i.e., including all flows. Per flow records are not available via this API.

The statistics reported in response to this command comply to the corresponding IANA parameters. The statistic counts are the 'total' counts (accumulated since the flow group is created) and are not incremental counts.

Table 55: Statistics Object Parameters

Parameter	Type	Description
name	String	Name of the flow group
octetTotalCount	Integer	Total number of bytes that belonged to the flow group.
packetTotalCount	Integer	Total number of packets that belonged to the flow group.
tcpSynTotalCount	Integer	Total number of TCP SYN messages that belonged to the flow group.
tcpFinTotalCount	Integer	Total number of TCP FIN messages that belonged to the flow group.
tcpRstTotalCount	Integer	Total number of TCP RST messages that belonged to the flow group.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-flowgroup-statistics",
  "asic-id": "1",
  "params": {
    "name": "httpsrv1"
  },
  "id": 1
}
```

The Agent acknowledges the command and returns an appropriate error code if needed. For an immediate response, the Agent returns the desired statistics.

A sample response is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-flowgroup-statistics",
  "asic-id": "1",
  "version": "4",
  "result": [{
    "name": "httpsrv1",
    "octetTotalCount": 12000,
    "packetTotalCount": 420,
    "tcpSynTotalCount": 25,
    "tcpFinTotalCount": 20,
    "tcpRstTotalCount": 2,
  }],
  "id": 1
}
```

8.5.3.4 get-ft-capabilities

The `get-ft-capabilities` API retrieves the current capabilities and the scaling parameters of the Agent. There are no parameters associated with the API.

The response is a list of parameters that are described below.

Table 56: get-ft-capabilities Response Parameters

Parameter	Type	Description
<code>multiple-collectors</code>	Boolean	Indicates whether multiple collectors are supported by the Agent.
<code>max-collectors</code>	Integers	Maximum number of collectors that can be supported. Range is 1 - 16
<code>collector-load-balancing</code>	Boolean	Indicates whether load balancing across various collectors is supported by the Agent. This parameter may not be part of the response if the multiple collectors are not supported.
<code>export-protocols</code>	Array of Strings	An array of strings that indicate the export protocols that are supported. Valid values are: "IPFIXv10"
<code>flowgroup-types</code>	Array of Strings	An array of strings that indicate the types of flow groups that are supported. Valid values are: "five-tuple", "egress-ports", "ingress-ports", "congestion"
<code>max-flow-groups</code>	Integer	Maximum number of flow groups that can be supported by the Agent
<code>max-flows</code>	Integer	Maximum number of flows that can be monitored simultaneously
<code>flow-sampling</code>	Boolean	Indicates whether flows can be sampled and reported to BroadView Client.
<code>flow-sampling-mode</code>	String	If the flow-sampling is supported, this parameter indicates how the flows are sampled. Valid values are: "CPU", "TCB" "TCB" indicates that the sampling is supported in hardware and therefore more efficient and accurate. "CPU" indicates that the sampling is done using CPU intervention. This may not be efficient. This parameter may not be part of the response if the flow-sampling is not supported.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-capabilities",
  "asic-id": "1",
  "params": { },
  "id": 1
}
```

A sample response is provided below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-capabilities",
  "asic-id": "1",
  "version": "4",
  "result": {
    "multiple-collectors": 0,
    "max-collectors": 1,
    "export-protocols": ["IPFIXv10"],
    "flowgroup-types": ["five-tuple"],
    "max-flow-groups": 255,
    "max-flows": 25000,
    "flow-sampling": 0
  },
  "id": 1
}
```

8.5.3.5 get-ft-status

The `get-ft-status` API retrieves the status of the Flow Tracker feature. There are no parameters associated with the API.

The parameters that are listed for the `configure-ft-feature` API are returned in response to this API. In addition, the following parameters are also included in the response.

Table 57: get-ft-status Response Parameter

Parameter	Type	Description
num-flow-groups	Integer	Number of flow groups currently configured on the Agent.
flow-group-list	Array of Strings	Array of all the names of the flow groups created on the Agent so far.
num-collectors	Integer	Number of collectors currently configured on the Agent.
collector-list	Array of Strings	Array of all the names of the collectors created on the Agent so far.

A sample JSON-RPC request is shown below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-status",
  "asic-id": "1",
  "params": { },
  "id": 1
}
```

A sample response is provided below.

```
{
  "jsonrpc": "2.0",
  "method": "get-ft-status",
  "asic-id": "1",
  "version": "4",
  "result": {
    "ft-enable": 1,
    "ipfix": {
      "template-update-periodicity": 600,
      "flow-records-src-port": 8085,
      "flow-group-records-src-port": 8086,
      "flow-record-observation-domain": 1,
      "flow-group-record-observation-domain": 2
    }
  },
  "num-flow-groups": 1,
  "flow-group-list": ["httpsrv1"],
  "num-collectors": 1,
  "collector-list": ["ntop-srv1"]
},
  "id": 1
}
```

8.6 Configuration Persistence

All the Flow Tracker configuration received via the REST API must be persistently stored and reapplied upon a restart of the Agent application. The only exception is the automatic periodic reporting of the flow group statistics, which needs to be explicitly re-enabled by the client.

Chapter 9: Reference Applications

Also available, along with the Agent, is the BroadView Instrumentation library. This library is Java-based, runs on a PC, and allows clients to interact with multiple switches running the BroadView Agent. This application showcases the visibility and telemetry information provided by the Agent.

Chapter 10: License

All of the user-space code of BroadView is licensed under the Apache License, Version 2.0 (the *License*). You can obtain a copy of this license at <http://www.apache.org/licenses/LICENSE-2.0>.

The sFlow feature of the software is an implementation of the specifications published at <http://www.sflow.org> and governed by the sFlow License Agreement: <http://www.inmon.com/technology/sflowlicense.txt>.

It may be noted that the actual sFlow functionality is provided by the silicon. The Agent enables this functionality upon a client request. Consult the user documentation of the silicon for the details, such as how sample packets are packaged into sFlow frames, packet formats, and so on.

Chapter 11: Source Code

BroadView software is available in two packages¹⁵:

- An OEM and ODM Development Package (ODP), which is a full source code package distributed under Broadcom SLA.
- A Community Development Package (CDP), which distributed via GitHub.

15. Not all packages may be available in a given release. Consult the release documentation for specific packages available with a given Agent release.

Chapter 12: Documentation

Technical documents are located in the release package, including:

- Compile/Build procedure
- Doxygen-generated documentation that discusses development of custom SB plug-ins

Appendix A: JSON Payload Revision History

Table 58: JSON Payload Revision History

Previous Version	New Version	Change Description	Example
0	1	New fields introduced in response message for the <code>get-bst-feature</code> command. <ul style="list-style-type: none"> stats-in-percentage async-full-reports trigger-rate-limit trigger-rate-limit-interval send-snapshot-on-trigger 	<pre>{ ... "method": "get-bst-feature", "result": { "stat-in-percentage": 0, "async-full-reports": 1, "trigger-rate-limit": 0, "trigger-rate-limit-interval": 0, "send-snapshot-on-trigger": 0 }, ... }</pre>
0	1	New fields introduced in the notification message of <code>trigger-report</code> : <ul style="list-style-type: none"> realm counter 	<pre>{ ... "method": "trigger-report", "realm": "ingress-port-priority-group", "counter": "um-share-buffer-count", ... }</pre>
1	2	New fields introduced in the response message for the <code>get-switch-properties</code> command: <ul style="list-style-type: none"> uid agent-ip agent-port time-stamp 	<pre>{ ... "method": "get-switch-properties", "time-stamp": "2015-10-18 - 00:15:04", "uid": "0000d80bb99bbbbb", "agent-ip": "192.168.1.2", "agent-port": "8080" ... }</pre>
2	3	New fields introduced in the response/notification messages for the following commands: <ul style="list-style-type: none"> get-packet-trace-lag-resolution, get-packet-trace-ecmp-resolution get-packet-trace-profile packet-info packet-received-time-stamp packet-ingress-time-stamp packet-egress-time-stamp packet 	<pre>{ ... "time-stamp": "2014-11-18 - 00:15:04 ", "packet-info": [{ "packet-received-time-stamp": "2014-11-18 - 00:15:00", "packet-ingress-time-stamp": "2014-11-18 - 00:15:00", "packet-egress-time-stamp": "2014-11-18 - 00:15:00", "packet": "0010203232.." }], ... }</pre>

Table 58: JSON Payload Revision History (Continued)

Previous Version	New Version	Change Description	Example
3	4	New API <code>get-port-queue-map</code> in BST. New fields introduced in the API/response/notification messages to support hardware-queue to user-queue mapping.	<pre>{ "jsonrpc": "2.0", "method": "configure-bst-thresholds", "asic-id": "1", "params": { "realm": "egress-uc-queue", "port": "2", "user-queue": 0, "uc-threshold": 1000 }, "id": 1 }</pre>
4	5	New fields introduced in the response message for the <code>get-switch-properties</code> command: <ul style="list-style-type: none"> ■ <code>last-saved-configuration</code> ■ <code>agent-up-time</code> 	<pre>{ ... "method": "get-switch-properties", "time-stamp": "1970-01-02 - 09:40:43", "last-saved-configuration": "1969-12-31 - 16:00:00", "agent-up-time": 4 ... }</pre>
5	6	New APIs introduced in Packet Trace feature <ul style="list-style-type: none"> ■ <code>get-packet-drop-reason</code> ■ <code>get-packet-dropctrs</code> Additional field introduced in heartbeat messaging indicating unsaved configuration.	<pre>{ ... "method": "get-packet-drop-reason", "params": { "packet": "000001000002...", "port": "1" }, ... }</pre>

Table 58: JSON Payload Revision History (Continued)

Previous Version	New Version	Change Description	Example
6	7	<p>New APIs supported in the Flow Tracker feature</p> <ul style="list-style-type: none"> ■ remove-ft-collector ■ clear-ft-flowgroup-statistics ■ remove-ft-flowgroup ■ get-ft-collector <p>Additional field introduced in configure-ft-feature and create-ft-flowgroup and get-ft-capabilities response is enhanced to reflect the "egress-ports" flowgroup-types</p>	<pre>{ ... "method": "get-ft-capabilities", "version": "7", "result": { "multiple-collectors": 0, "max-collectors": 1, "export-protocols": ["IPFIXv10"], "flowgroup-types": ["five-tuple", "egress-ports"], "max-flow-groups": 255, "max-flows": 6144, "flow-sampling": 0 }, ... }</pre>

Revision History

BroadView-SP101; May 31, 2018

- Updated [Table 1, Supported Features by Silicon](#).
- Updated [Figure 1, BroadView Agent Architecture](#)
- Updated [Table 2, Error Codes](#) to add HTTP 204 and HTTP 403 status codes.
- Updated [Table 3, Switch Properties Parameters](#) (Added last-saved-configuration, agent-up-time, and unsaved-config parameters)
- Updated [Section 4.4, Configurability](#)
- Added [Section 4.5.2.2, save-configuration](#)
- Added [Section 4.5.2.3, restore-configuration](#)
- Added [Section 4.5.3.1, get-configuration](#)
- Updated sample response in [Section 4.5.3.4, get-switch-properties](#)
- Added [Section 5.7.2.4, configure-bst-multi-thresholds](#)
- Updated [Section 6.2, Configurability](#) for Packet Trace.
- Added [Section 6.7, Deciphering Drop-Profile](#)
- Added [Section 6.8.3.5, get-packet-drop-reason](#)
- Added [Section 6.8.3.6, get-packet-dropctrs](#)
- Added [Chapter 8, Flow Tracker](#)
- Updated [Chapter 10, License](#) to add sFlow information.
- Updated [Chapter 12, Documentation](#).
- Updated [Appendix A, JSON Payload Revision History](#)

BroadView-SP100; April 26, 2017

Initial release.

